

# SOFTWARE ENGINEERING (Week-6)

USAMA MUSHARAF

*LECTURER (Department of Computer  
Science)*

*FAST-NUCES PESHAWAR*

# CONTENTS OF WEEK # 6


## **Software Quality Attributes**

- Performance
- Scalability
- Reliability

## **Application Architecture to System Architecture**

# Sample System

---



# Performance


---

# Performance

---

Measure of how fast or responsive a system is under

- A given workload
  - Backend data
  - Request Volume
- A given hardware
  - Kind
  - capacity



# Performance Problems

---

How to spot a performance problem? How does it look like?

*Every performance problem is the result of some queue building somewhere.*

- *Network socket queue, DB IO queue, OS run queue etc.*

- Reasons for queue build-up
  - Inefficient slow processing
  - Serial resource access
  - Limited resource capacity



# Performance Principles

---

## Efficiency

- Efficient Resource Utilization
  - IO- Network, Memory, Disk
  - CPU

## Efficient Logic

- Algorithms
- DB Queries

## Efficient Data Storage

- Data Structures
- DB Schema


## Caching

## Concurrency

- **Hardware**
- **Software**
  - Queuing
  - Coherence
- **Capacity**

# Performance Objectives

- Minimize Request-Response Latency
  - Latency is Measured in Time Units
  - Depends on
    - Wait/Idle Time
    - Processing Time
- Maximize Throughput
  - Throughput is Measured as Rate of Request processing
  - Depends on
    - Latency
    - Capacity





# Performance Measurement Metrics

## Latency

- Affects- User Experience
- Desired- As low as possible

## Throughput

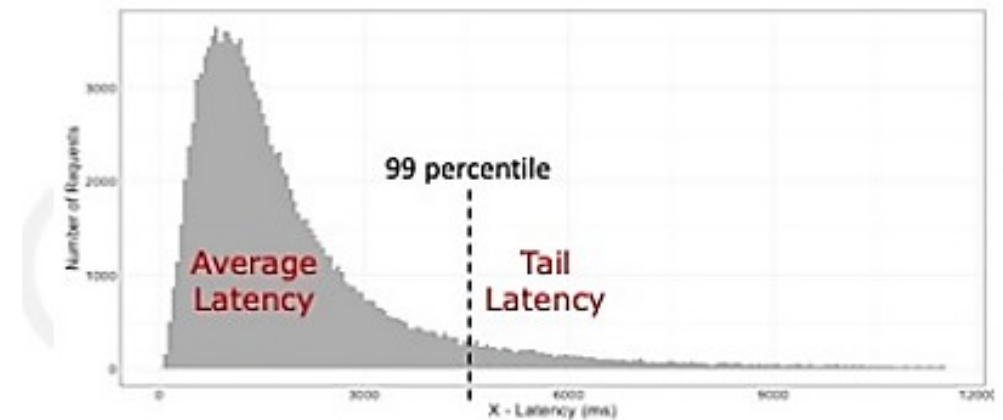
- Affects- Number of users that can be supported
- Desired- Greater than the request rate

## Errors

- Affects- Functional Correctness
- Desired- None

## Resource Saturation

- Affects- Hardware Capacity Required
- Desired- Efficient utilization of all system resources.

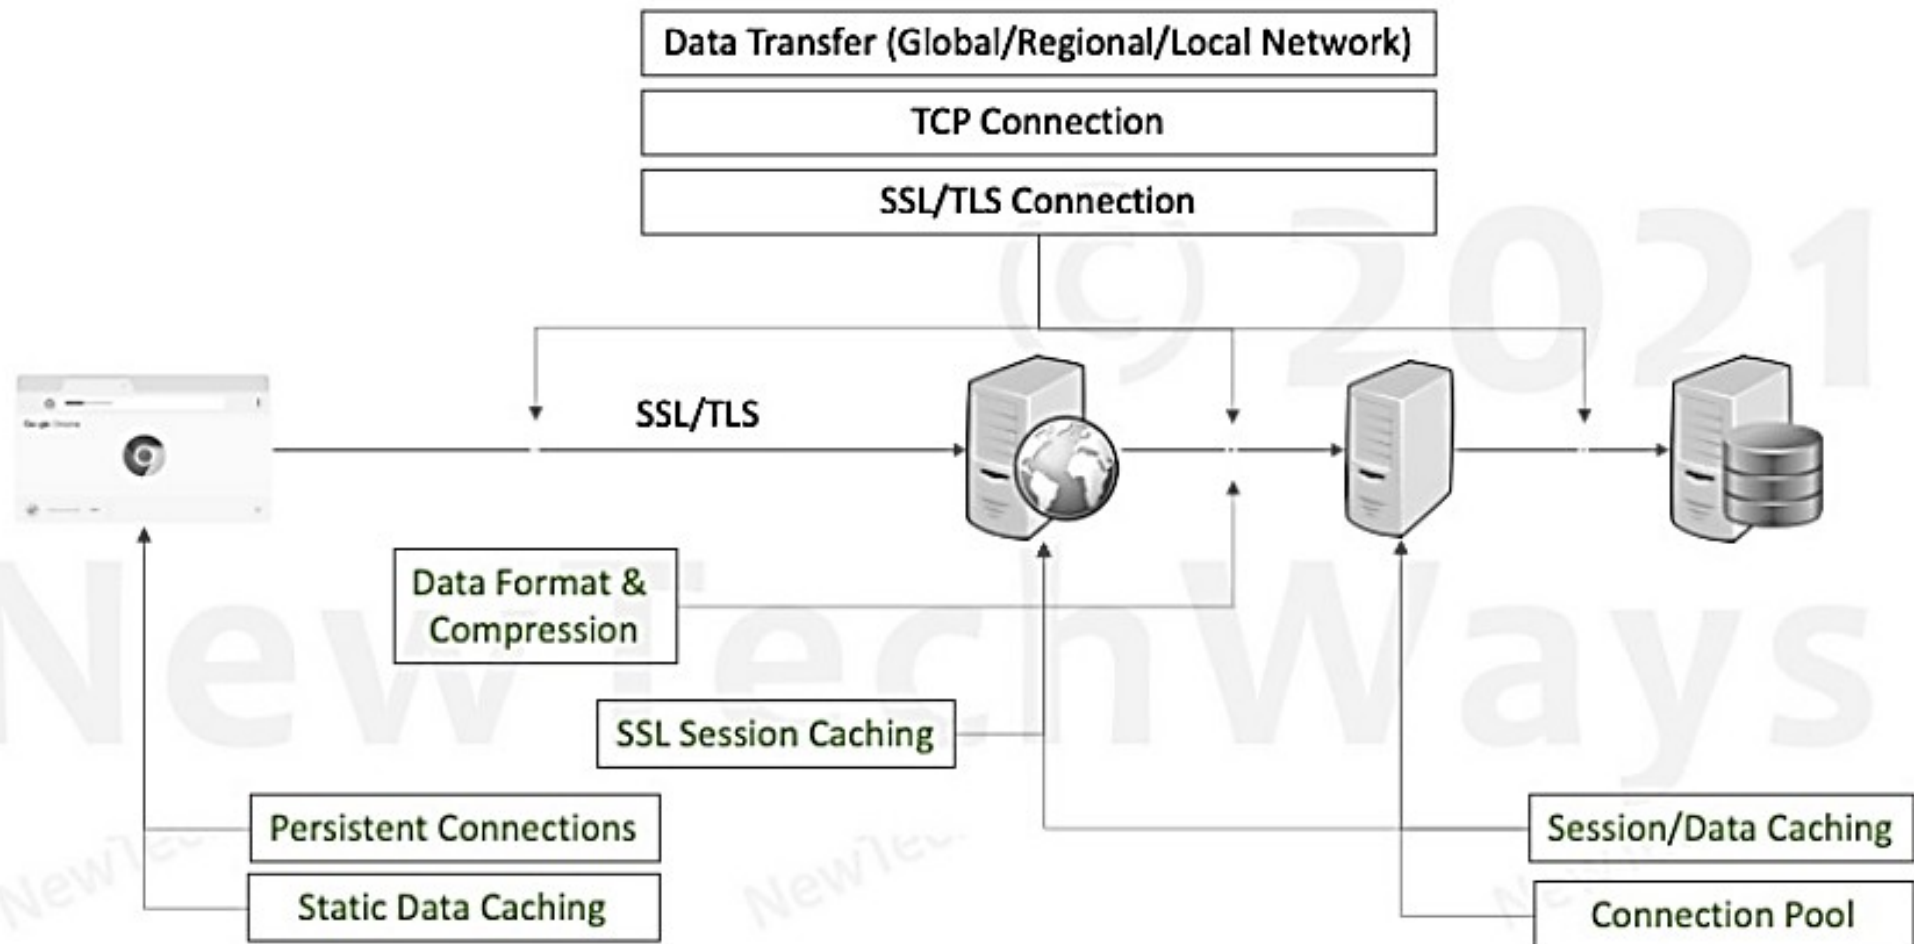


- Tail latency is an indication of queuing of requests
  - Gets worse with higher workloads
- Average latency hides the effects of tail latency
  - Also measure 99 (or 99.9) percentile latency

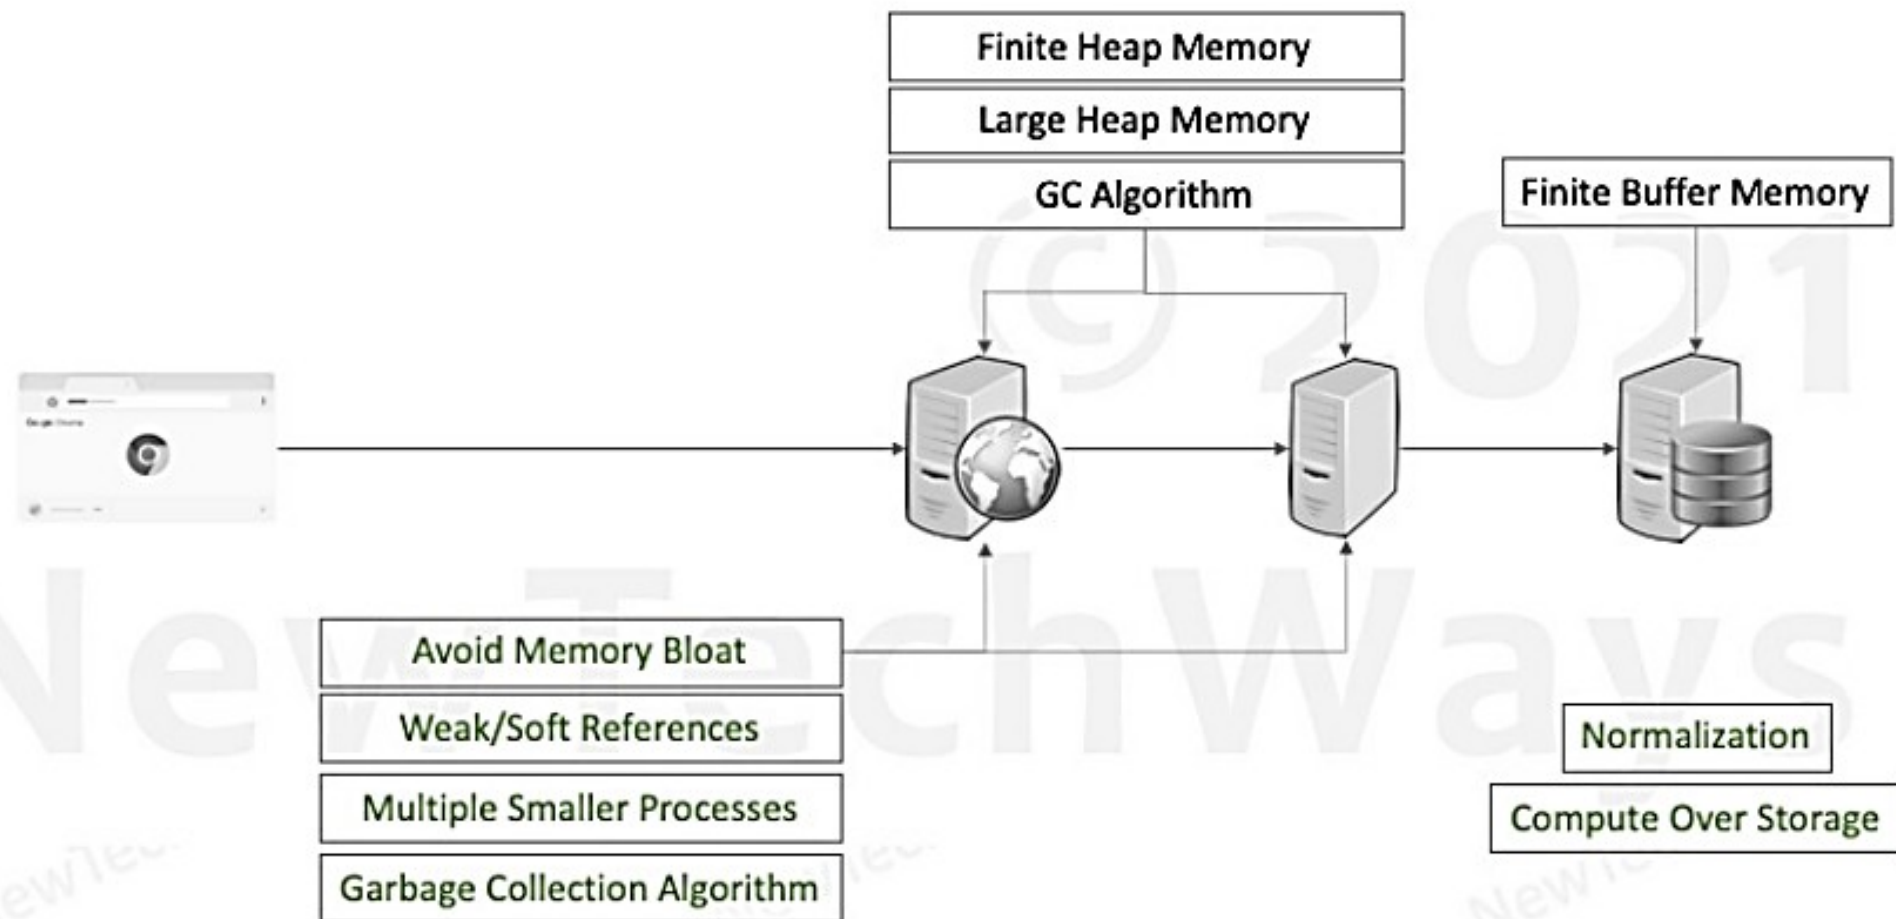
# Serial Request Latency

---

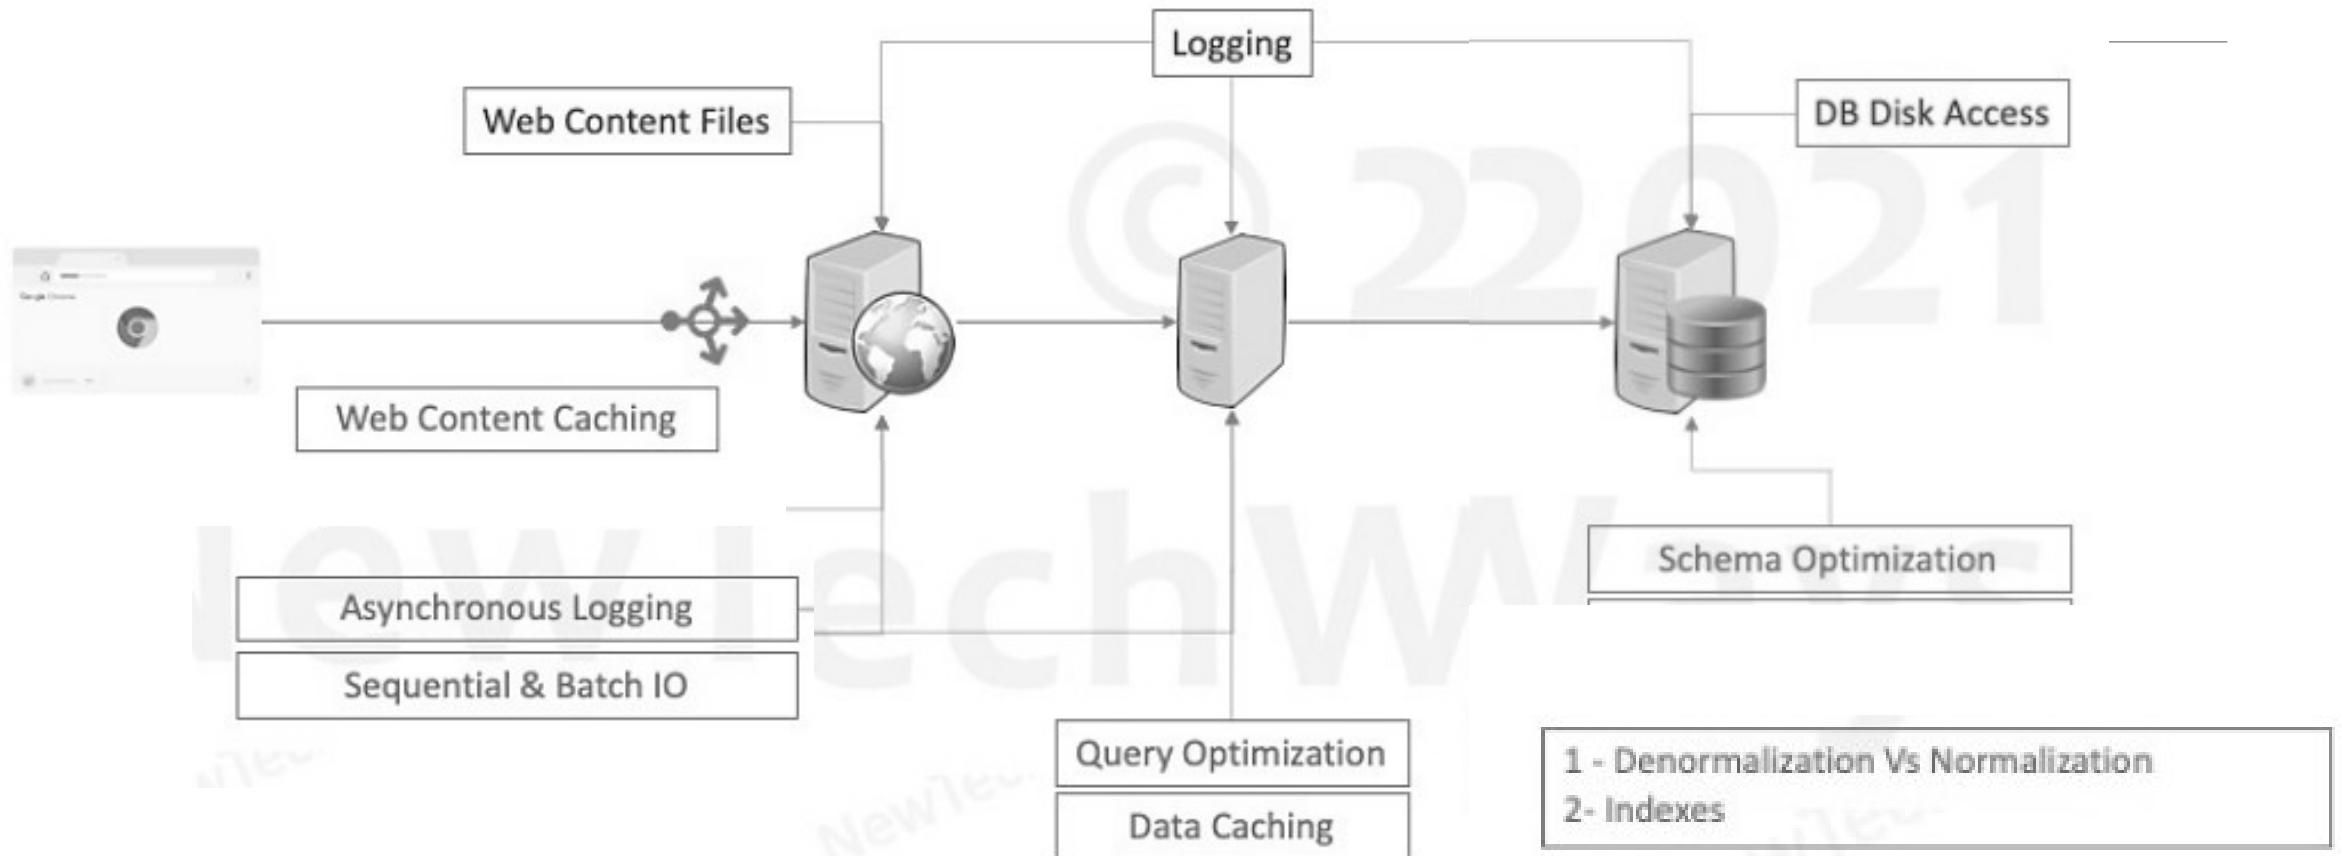
# Network Latency – Approaches



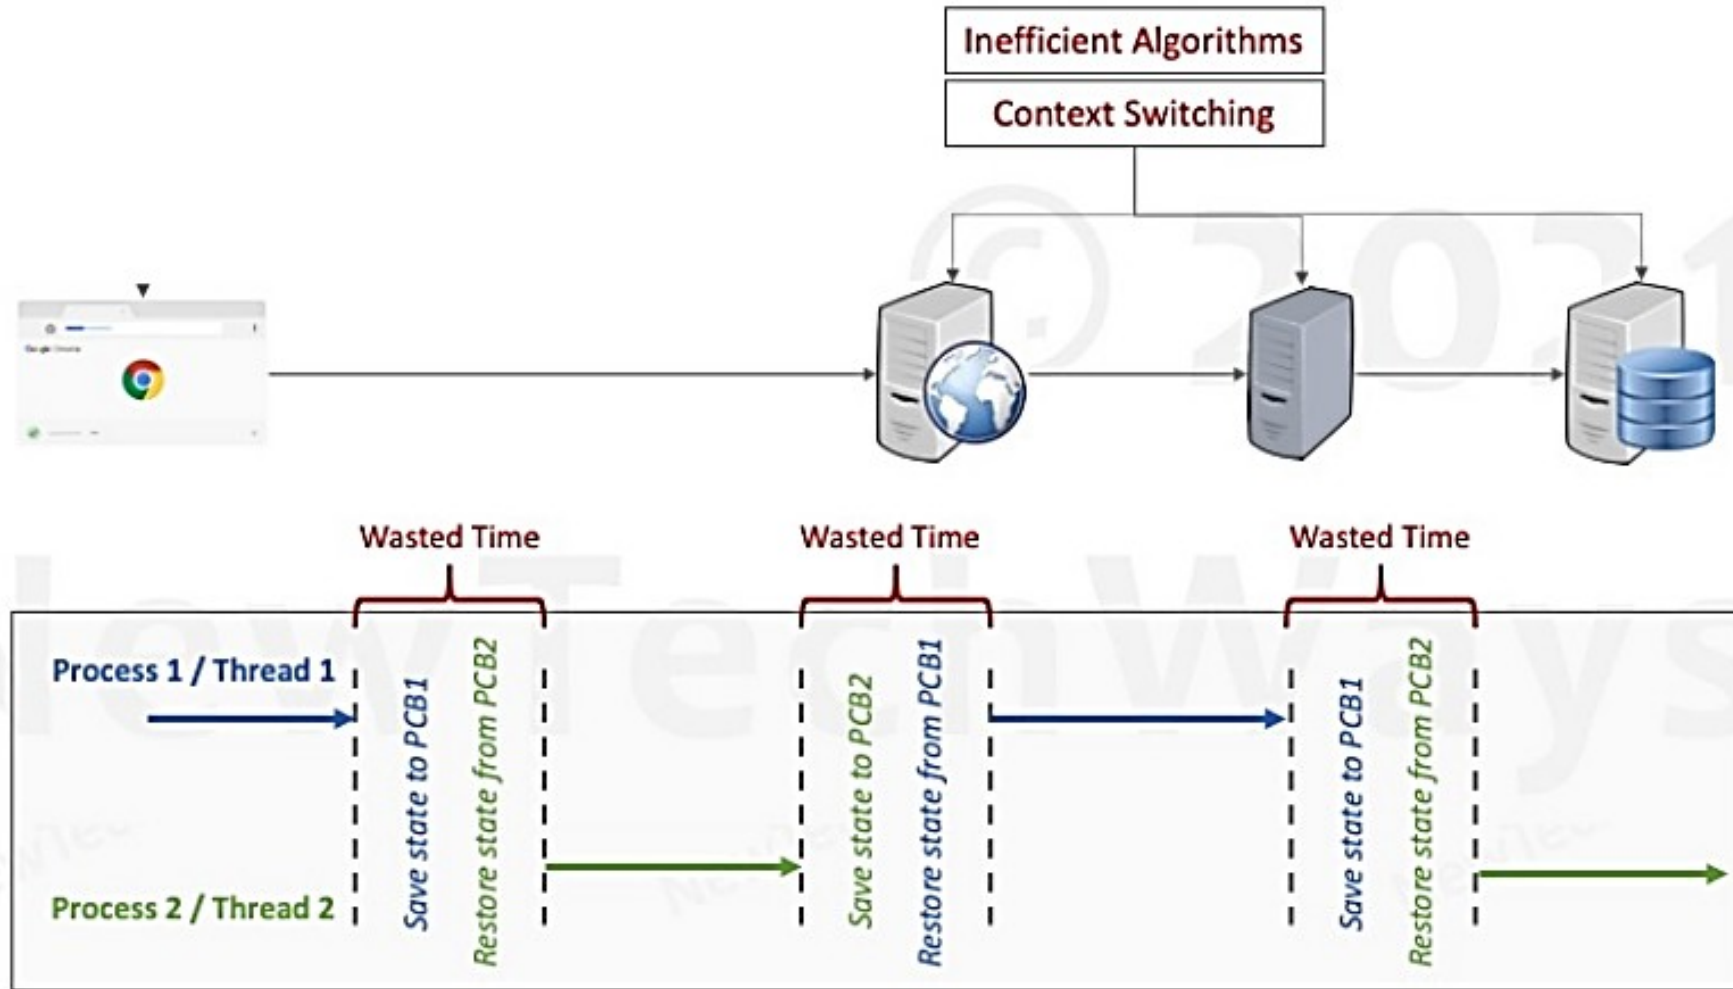
# Memory Latency – Approaches



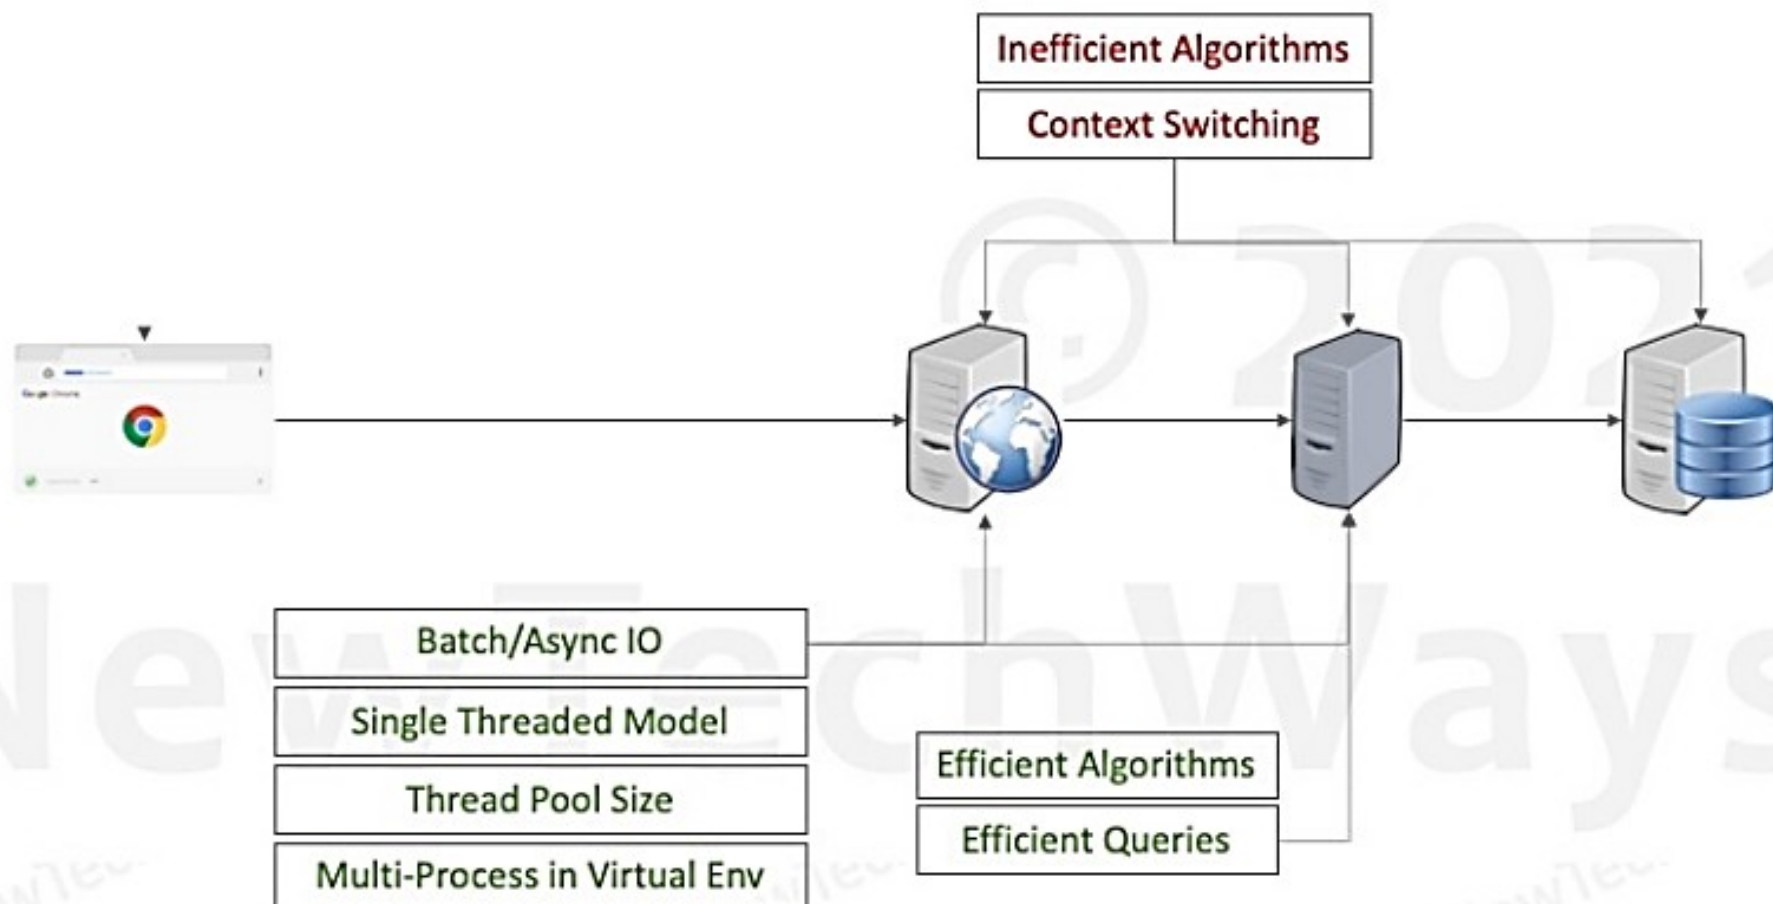
# Disk Latency - Approaches



# CPU Latency

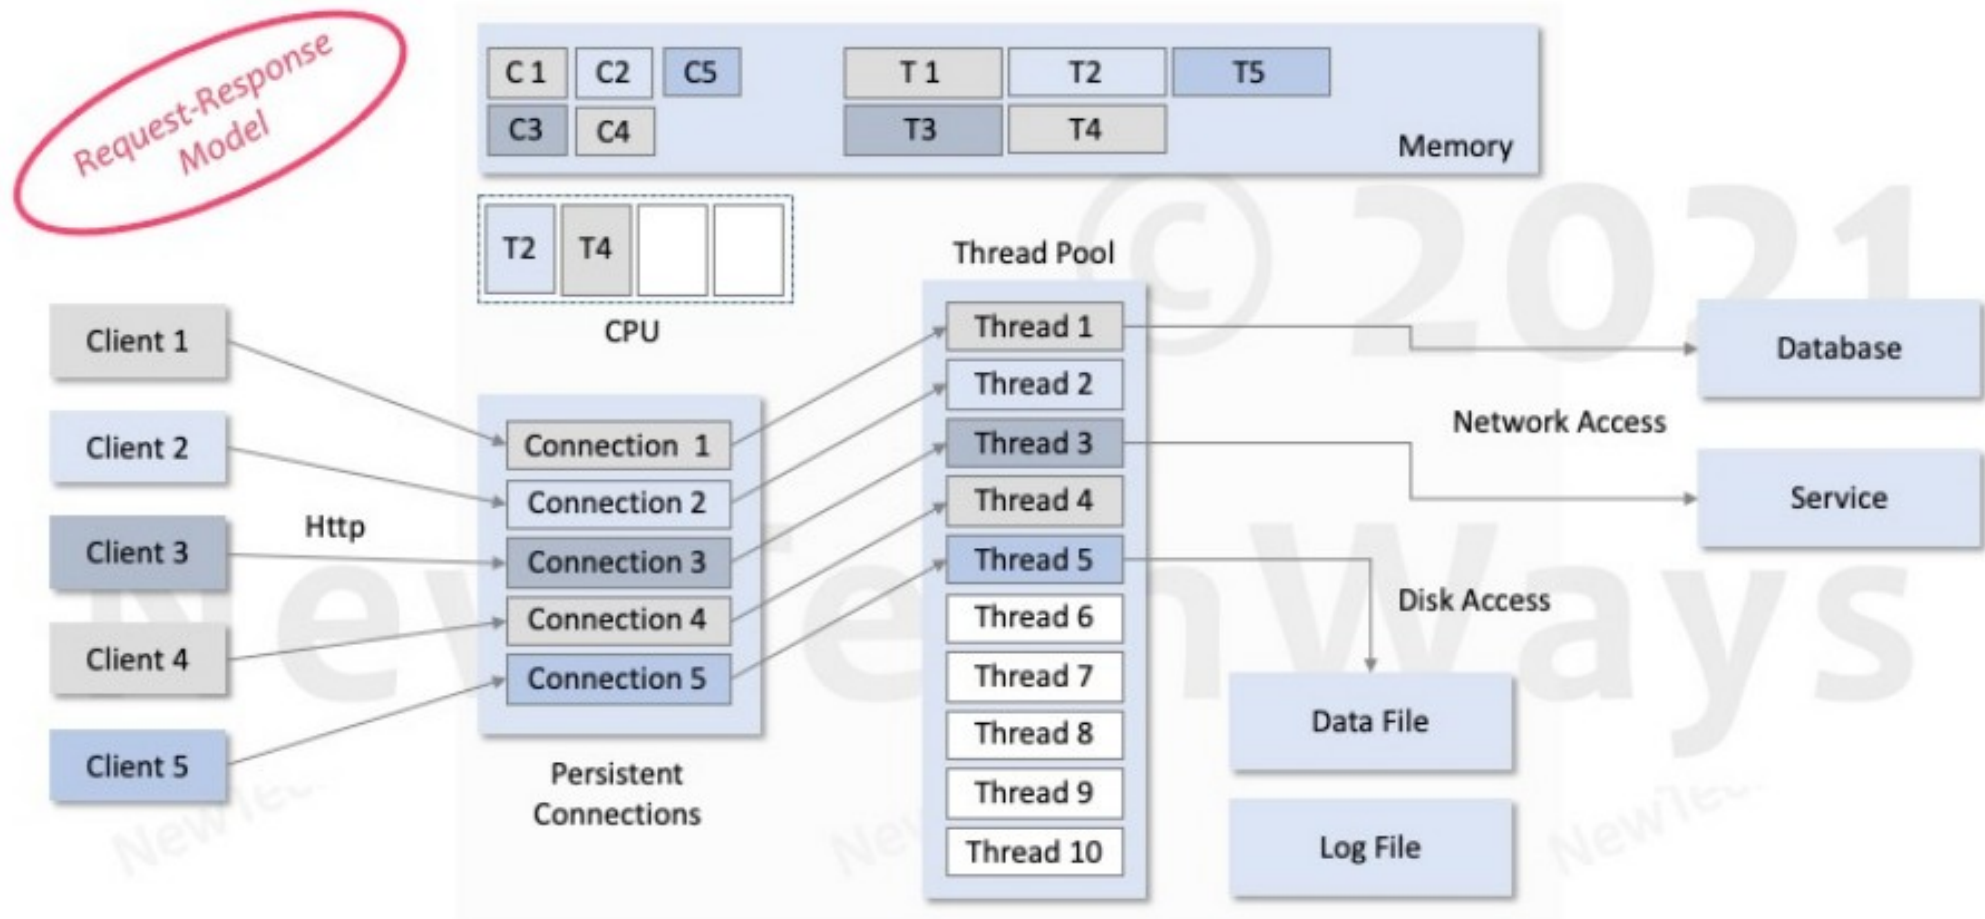


# CPU Latency – Approaches



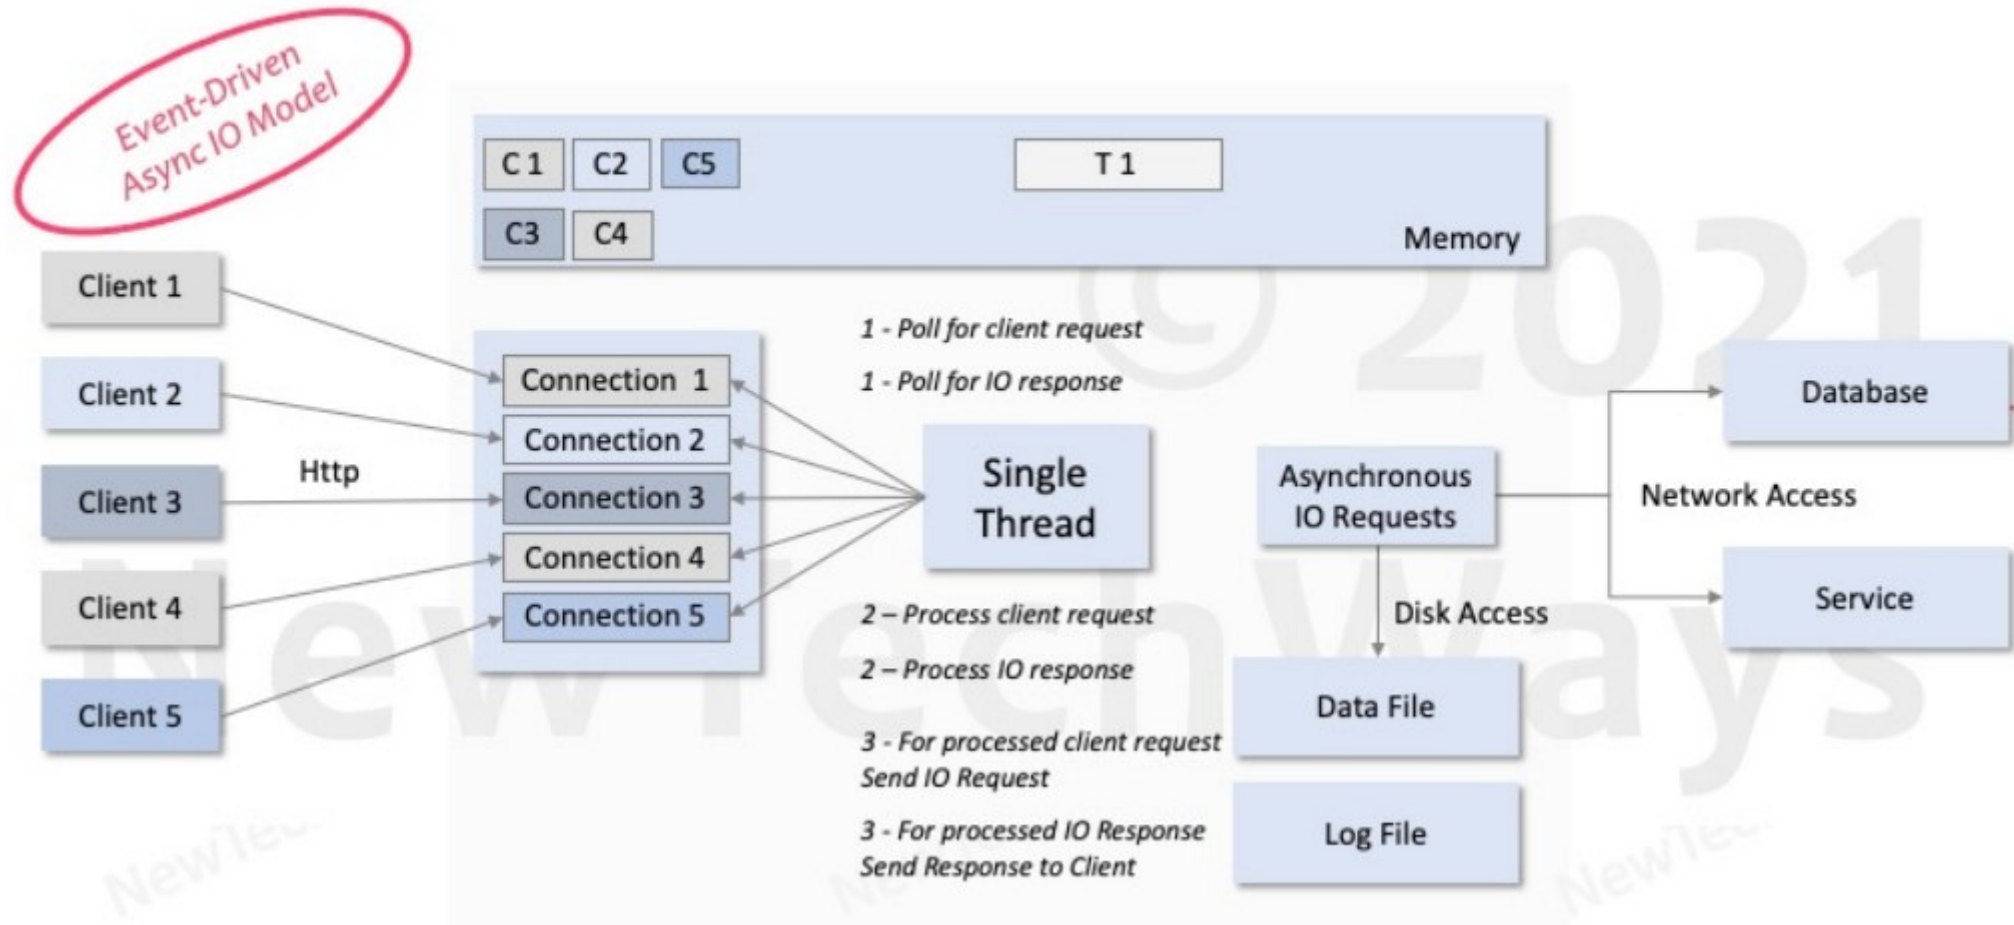


# Apache Webserver Architecture





# Nginx Architecture



# Parallel Request (Concurrency)

---



# Concurrent Processing

- Amdahl's Law

- $C(N) = \frac{N}{[1+\alpha(N-1)]}$

- C is capacity
- N is scaling dimension
  - like CPU or Load
- Alpha is resource contention
  - Alpha =0 , for linear performance

$C(N)$  = theoretical speed



Parallel + Serial

# Concurrent Processing

- Amdahl's Law


- $C(N) = \frac{N}{[1+\alpha(N-1)]}$

Queueing

- Universal Scalability Law

- $C(N) = \frac{N}{[1+\alpha(N-1)+\beta N(N-1)]}$

Queueing +  
Coherence




**C** is capacity

**N** is scaling dimension like CPU or load

**Alpha** represents resource contention

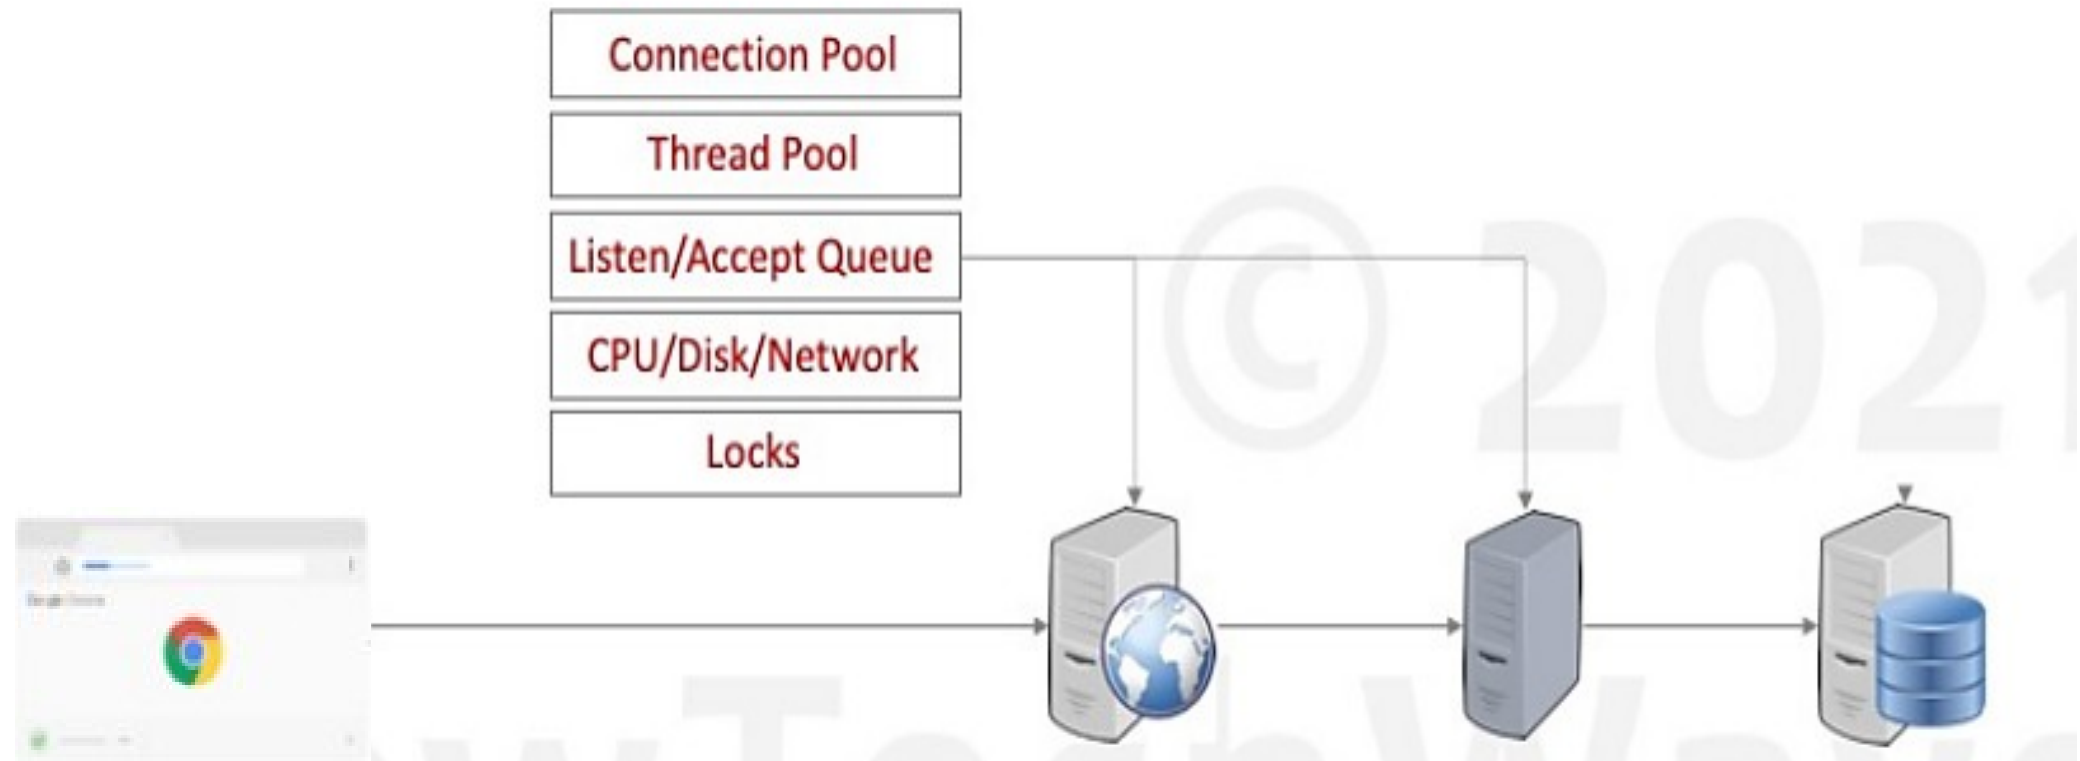
**Beta** represents coherence delay

Linear performance when alpha and beta are zero



# Contention

---



# Locking Compatibility Matrix

---

	S	X
S	true	false
X	false	false

# Coherence

---

Suppose that two threads are working on SharedObj. If two threads run on different processors each thread may have its own local copy of sharedVariable.

If one thread modifies its value the change might not reflect in the original one in the main memory instantly. This depends on the write policy of cache. Now the other thread is not aware of the modified value which leads to data inconsistency.

“Volatile” keyword in java tells the compiler that the value of a variable must never be cached as its value may change outside of the scope of the program itself.

# Scalability

---



# Performance vs Scalability

---

## FIXED LOAD

### Performance

- Low Latency
- High Throughput
  - Concurrency
    - Single-Machine- Multi-Threading
    - Multi-Machine- Multi-Threading + Multiprocessing = Distributed Processing

## VARIABLE LOAD

### Scalability


- High Throughput
  - Ability of a system to increase its throughput by adding more hardware capacity.
  - Both Ways: Up and Down

# Vertical & Horizontal Scalability

---


## Vertical

- Easier to achieve
- Limited scalability



## Horizontal

- Hard to achieve
- Unlimited scalability

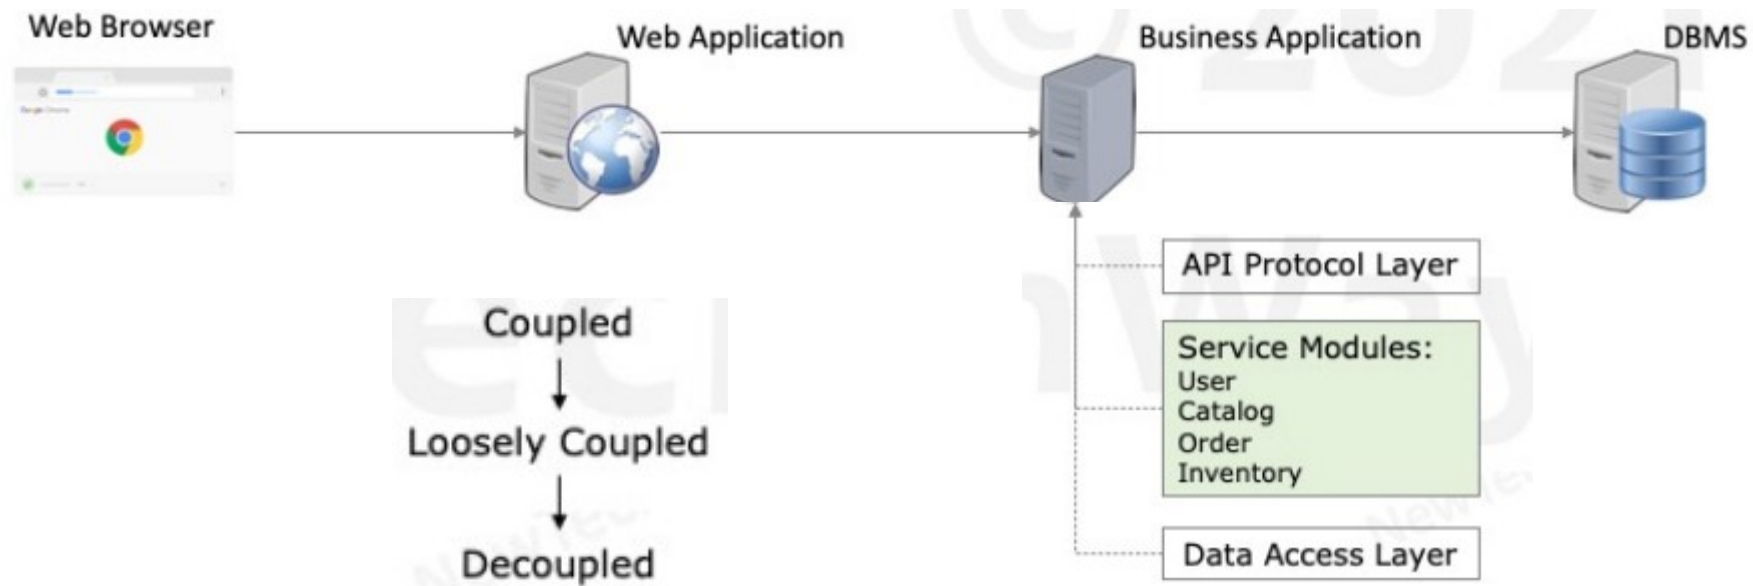


# Modularity

---

## Scalable architecture starts with Modularity

- Provides the foundation for breaking an application into more specialized functions/services.



# Horizontal Scaling Methods

---

## **1. Services**

## **2. Replication**

1. Stateful
2. Stateless

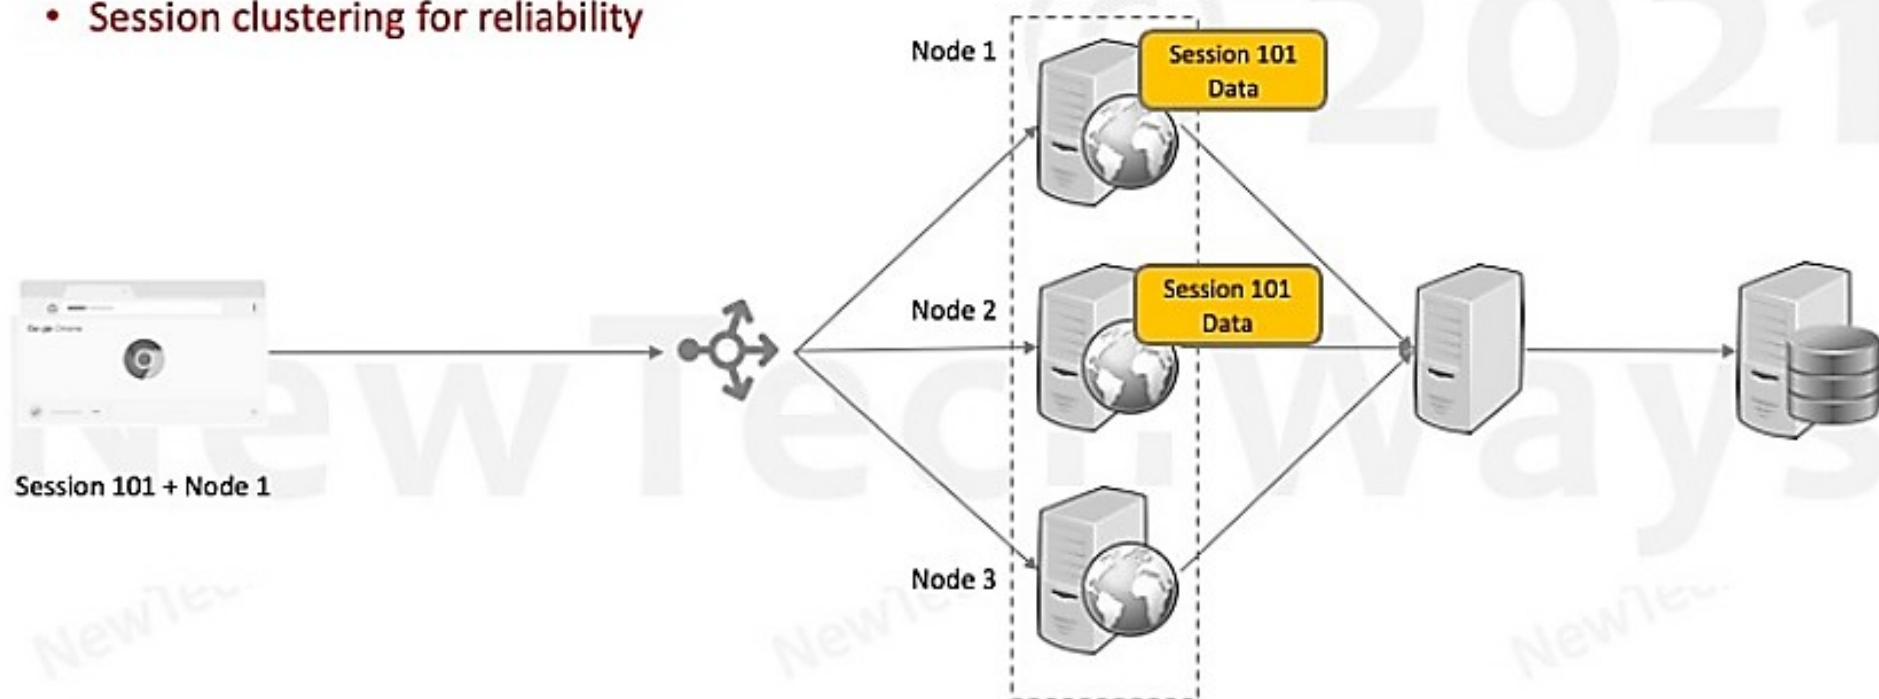
## **3. Portioning**

1. Vertical / Functionality Portioning
2. Database portioning

## **4. Asynchronous Calls**


# Web Stateful Replication

- When low latency is required
  - Sticky sessions/Session affinity
  - Sessions occupy memory
  - Session clustering for reliability



# Web Stateless Replication

- For higher scalability at the expense of higher latency
- Session data can be stored on
  - Client Side in Cookies
  - Server Side in Shared Cache



# Need For Specialized Services

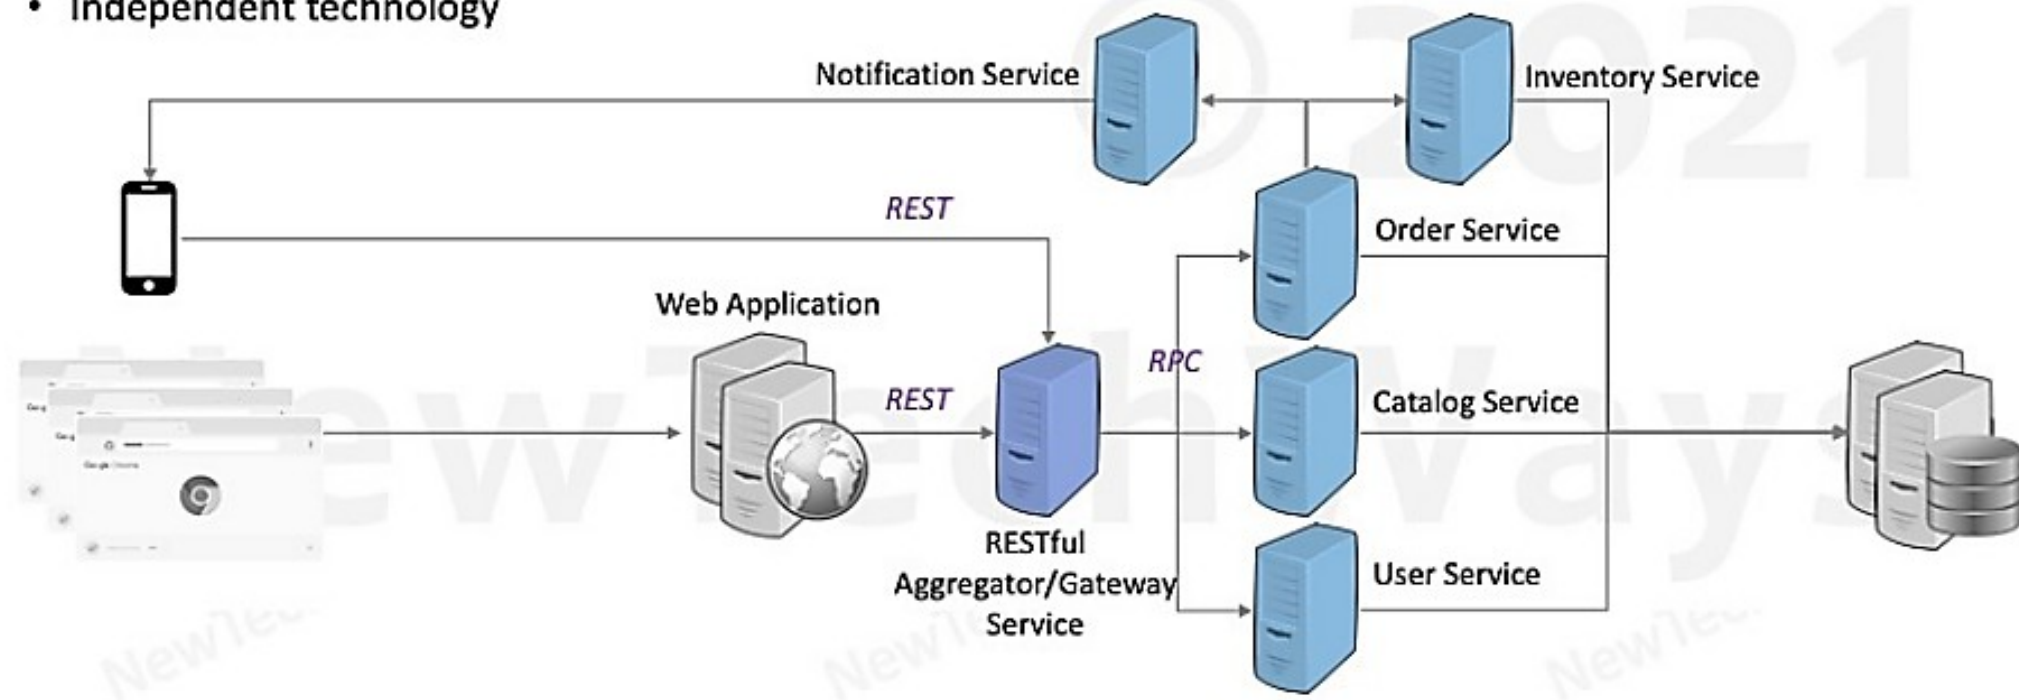
## Service Modules:

- User
- Catalog
- Order
- Inventory
- Notification



# Specialized Services (SOAP/REST-Services)


- Partially independent development and deployment
- Independent scalability
- Independent technology





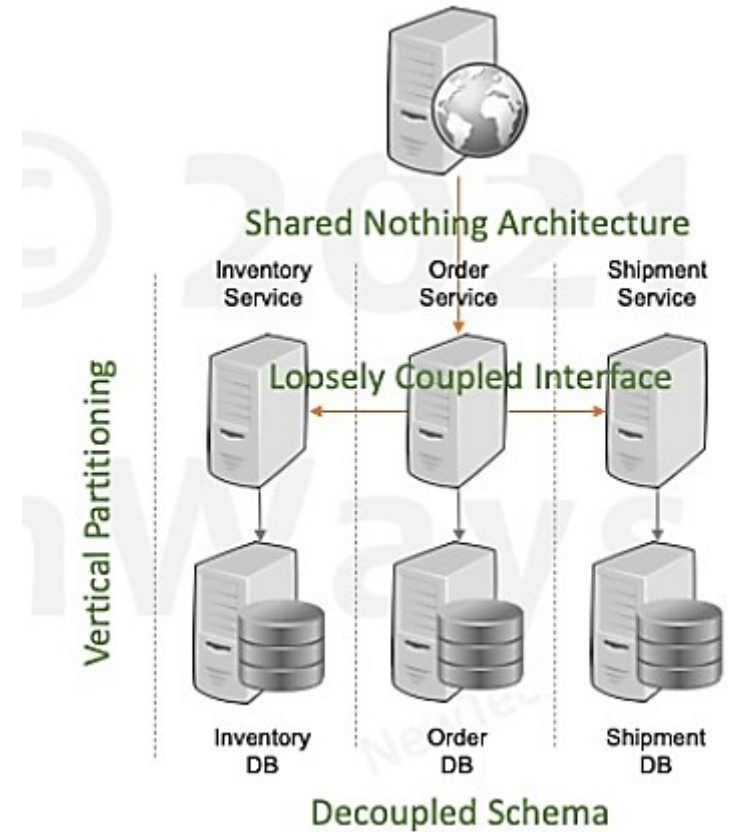
# Service Oriented Architecture

- Independent
  - Each service can have its own technology stack, libraries, frameworks etc.
  - Each service can be scaled independently and differently
- Not Independent
  - Common interface schema
    - XML schema
  - Common database schema
    - RDBMS schema
- Issues
  - Service development may be independent but not deployment
  - Single database has scalability limitations



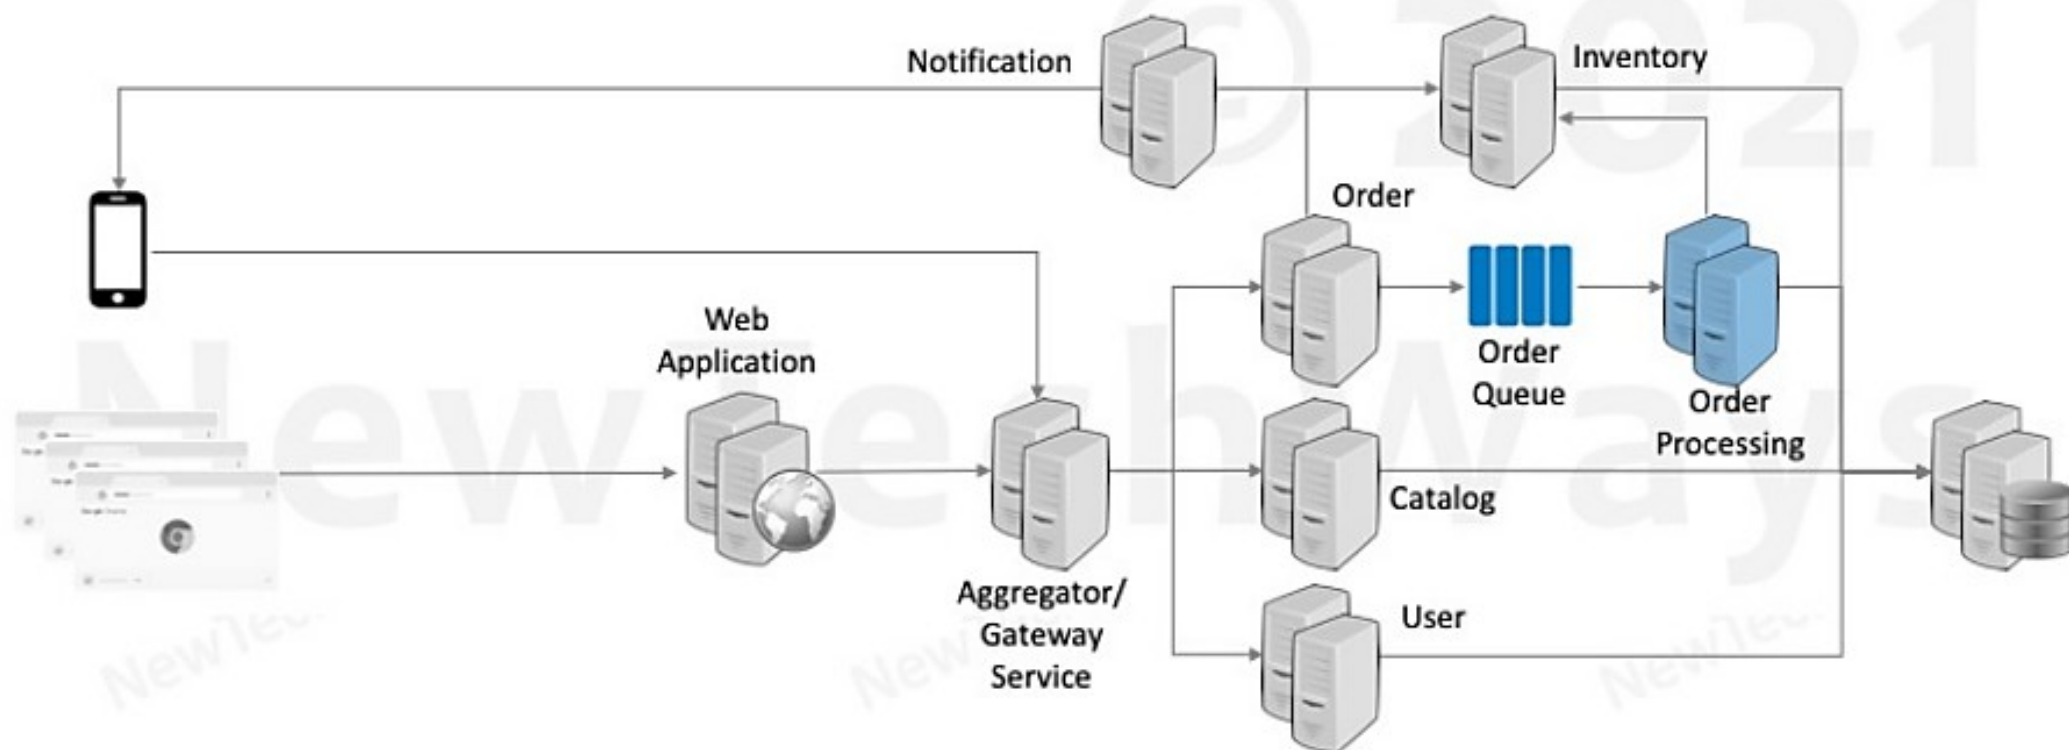
# Microservices

---




# Asynchronous Services

- Async services effectively reduces write load from a database

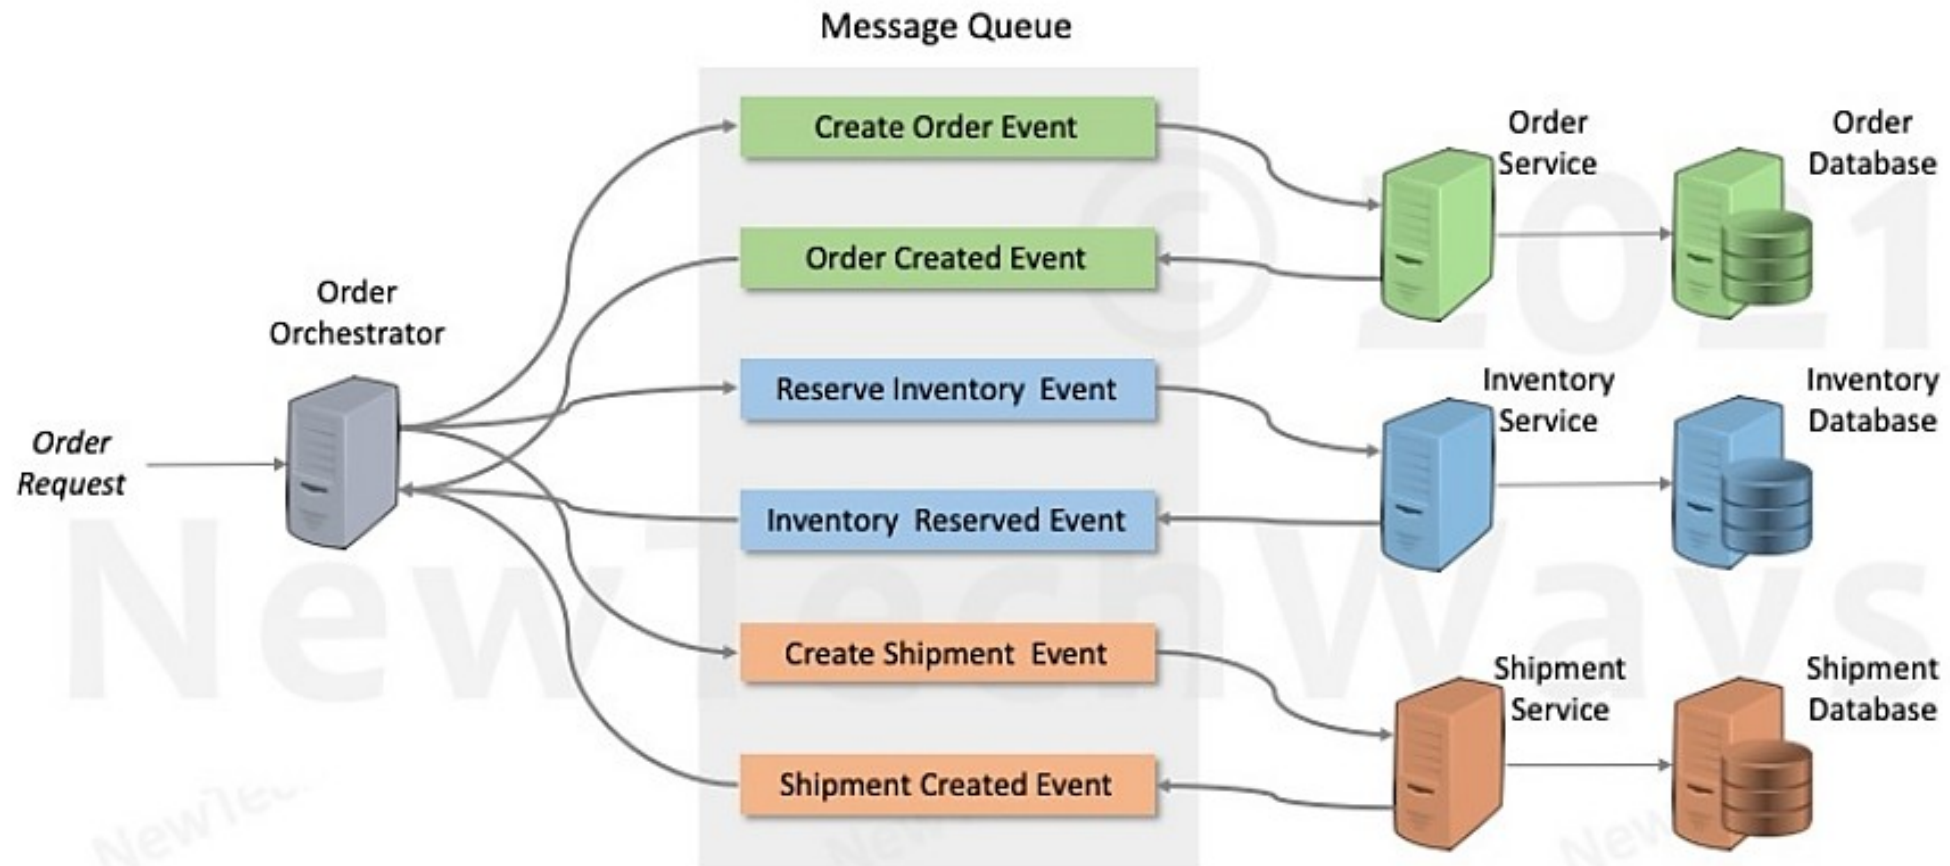


# Asynchronous Processing & Scalability

- Async services require infrastructure for average load as opposed to peak load



# Micro-Services Event Driven Transactions



# Reliability

---

# Reliability


---

Software reliability is defined as the probability of failure-free operation of a software system for a specified time in a specified environment.



# Failures in Large-Scale Systems

- Large scale systems are generally distributed systems
  - Large number of components
  - Large number of component instances
- Failures can be
  - Partial
  - Independent
  - Single point of failures
- Increased chance of partial failures
- Partial failures can lead to complete system failures





# Partial Failures

---

- Network Failure – LAN, WAN, Load Balancer
- Machine Failure – CPU, Disk, Memory
- Software Failure - Process
- Disaster – Datacenter
- Operations
  - Deployment Failure
  - Configuration Failure
  - Load Induced Failure
  - External Service Failure

- No matter how hard we try
  - Hardware and Networks will fail
  - A changing Software will fail
  - Disasters will happen

*After a point, its much more economical to recover from a failure instead of preventing it altogether*

# Reliability Engineering

---

1. Reliability
2. Availability
3. Fault Tolerance

# Availability

---

- It is the probability of a system working correctly at any given time and being available for operations

- Time based availability

$$\text{availability} = \frac{\text{uptime}}{(\text{uptime} + \text{downtime})}$$

- Request based availability

$$\text{availability} = \frac{\text{successful requests}}{\text{total requests}}$$

- There can be downtime but the system is expected to recover from the same in a quick time

# Fault Tolerance

---

- Fault Tolerance is a technique to improve Availability and/or Reliability of a system
- It is commonly referred to as an ability of a system to automatically
  - Detect partial failures
  - Handle partial failures
  - Recover from partial failures
- Serviceability
  - The ease with which a system can be serviced in the event of a failure also determines the availability of a system

# Fault Tolerant Design

---

1. Redundancy
2. Fault Detection
3. Recovery

# Redundancy


---

- Replication/Duplication of critical components or functions of a system in order to increase its reliability
- A secondary capacity is kept ready as a backup, over and above the primary capacity, in case the primary is not available

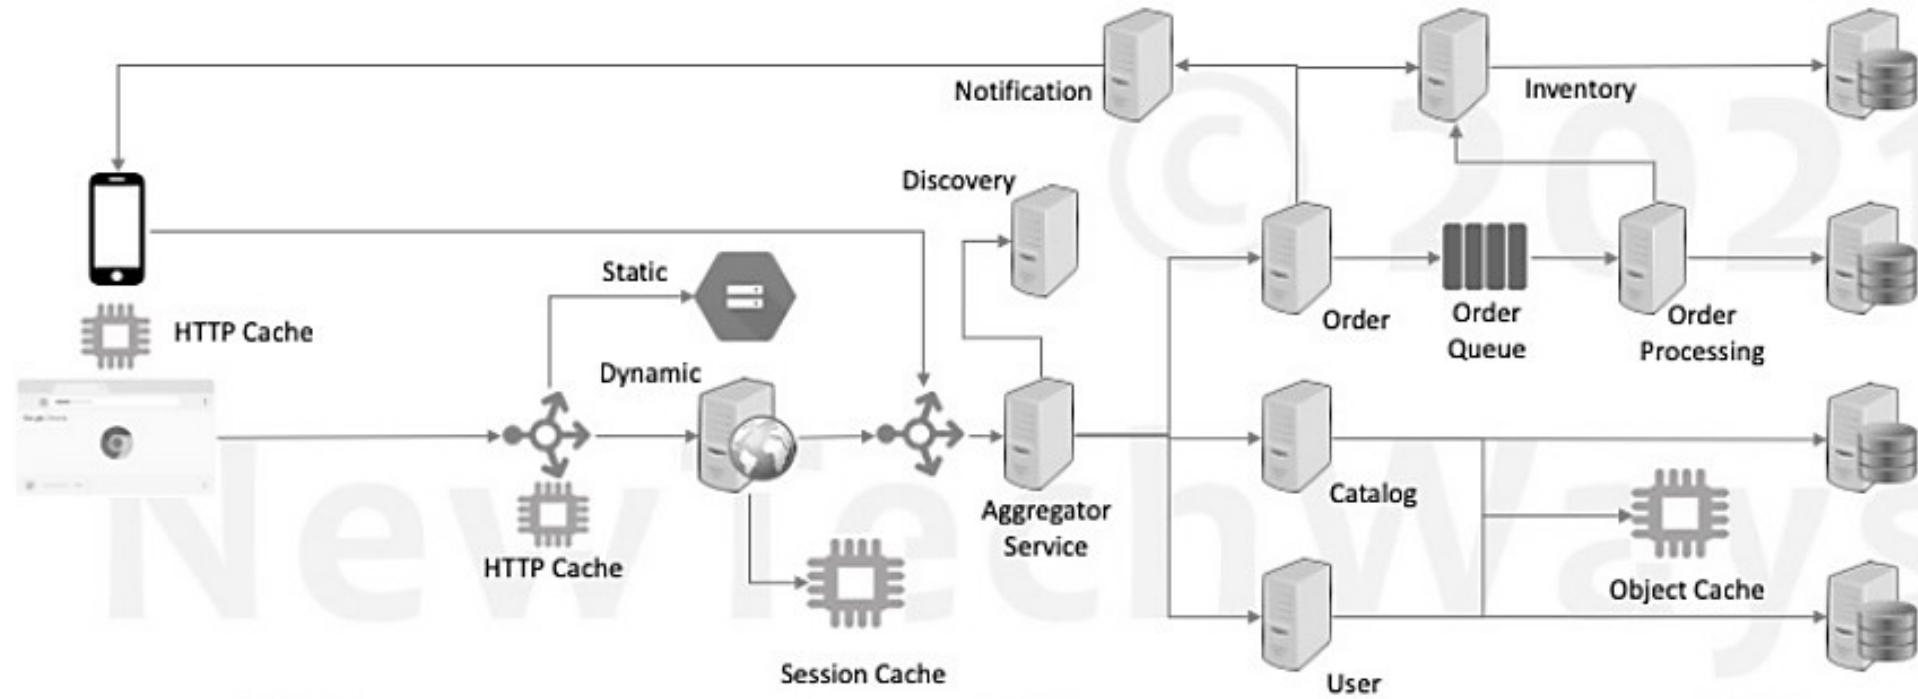


# Types of Redundancy

- **Active Redundancy – Hot Spare**
  - All nodes do the processing
  - Ideal for providing highest availability
- **Passive Redundancy – Warm Spare**
  - Only active nodes do the processing
  - Ideal for quick recovery
- **Cold Redundancy – Spare (Backup)**
  - Spare nodes are brought up only on a failover
  - It is not a high availability option

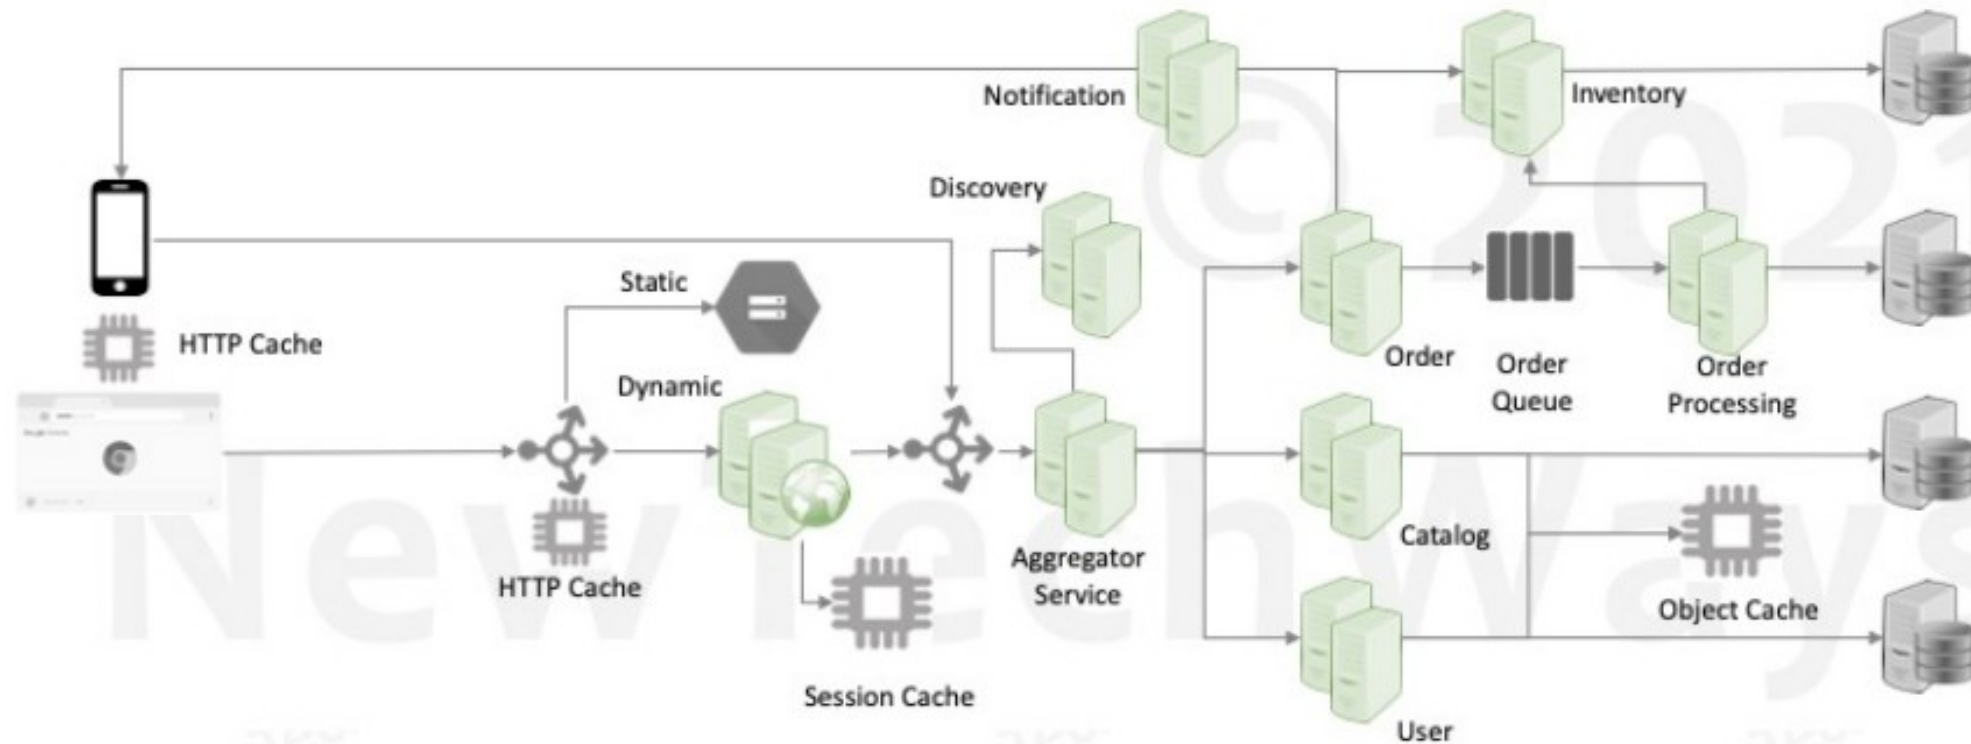


# Single Points of Failure

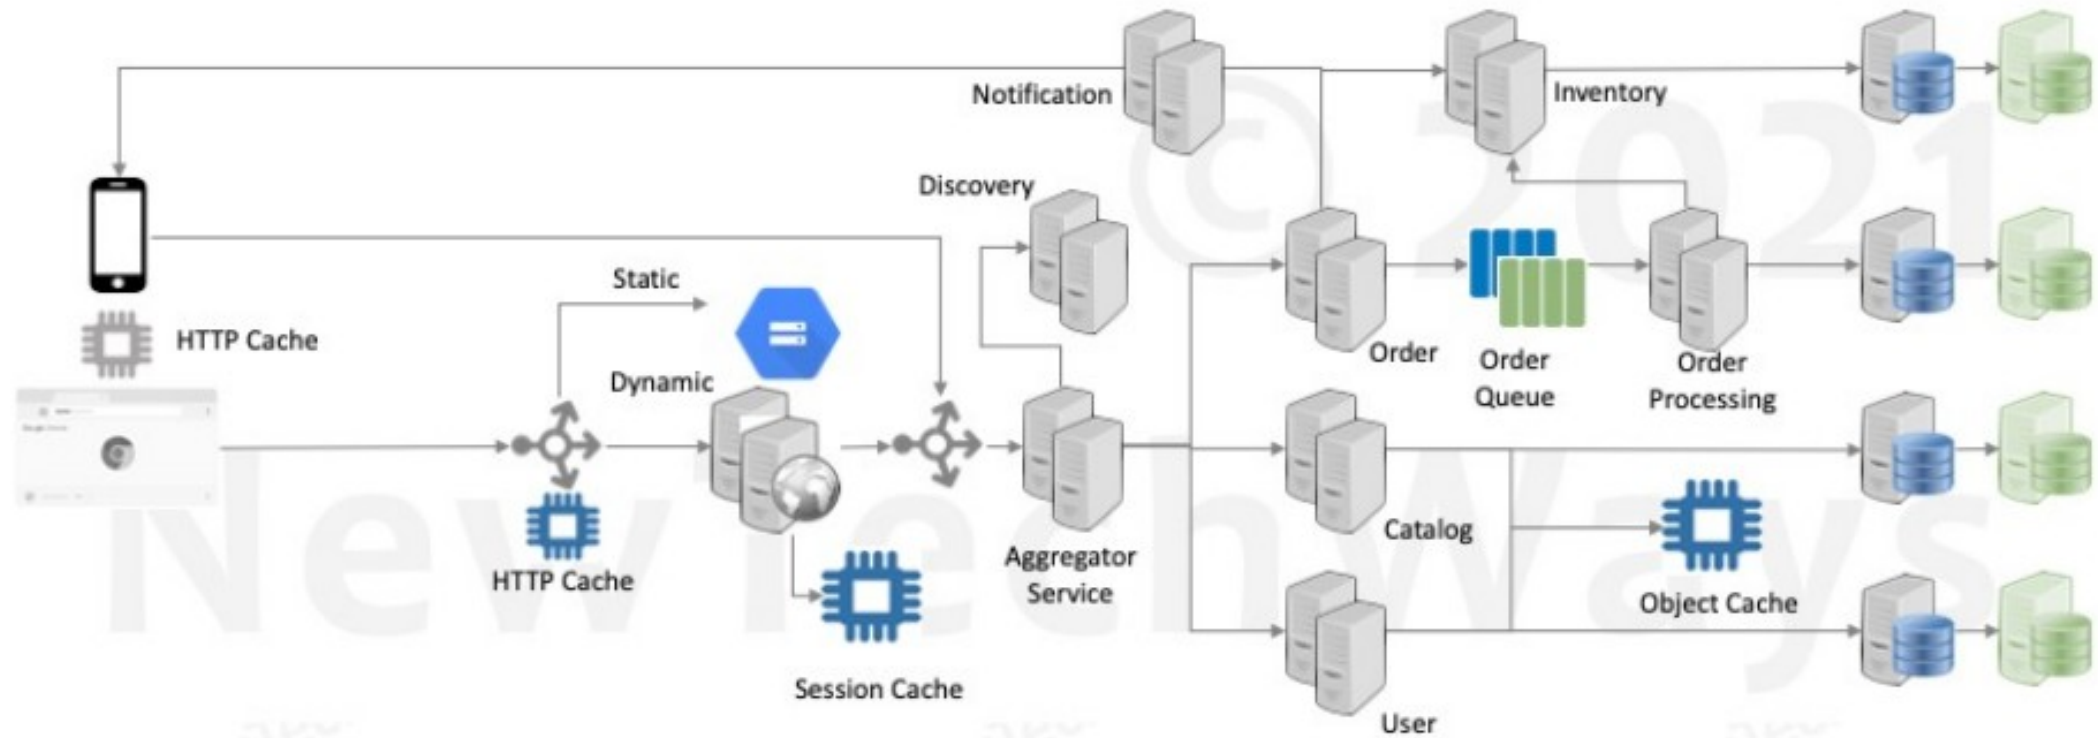




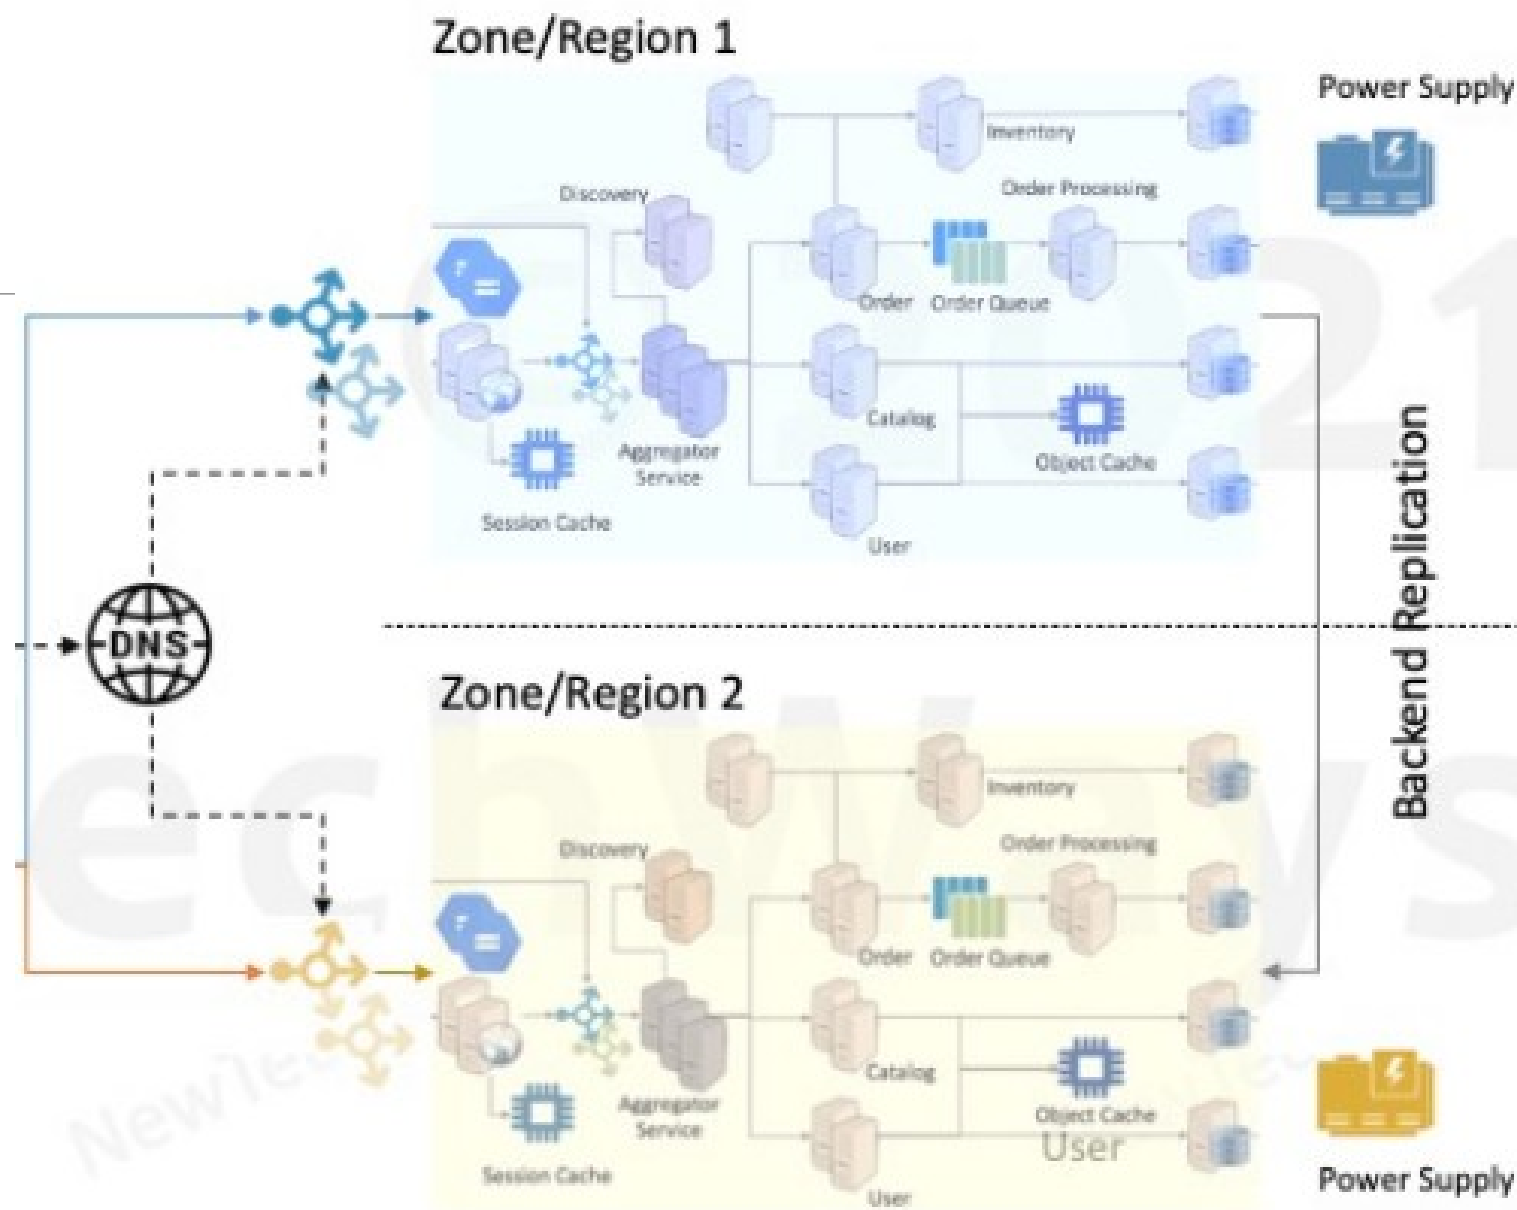
# Redundancy for Stateless Components



# Redundancy for Stateful Components



# Datacenter Redundancy



# Fault Models

---

- **Response Failure**
  - A server fails to receive or respond to incoming messages
- **Timeout Failure**
  - A server response duration is longer than timeout duration
- **Incorrect Response Failure**
  - A server's response is incorrect
- **Crash Failure**
  - A server halts but is working correctly until it halts
- **Arbitrary Response Failure**
  - A server's response is incorrect because its security is compromised

# Health Checks

---


## External Monitoring Service

- Ping based




## Internal Cluster Monitoring

- Heartbeat based





# External Monitoring Service

- Health check service generates
  - Alerts – for recovery
  - Events – for scaling
- Application Health Checks
  - HTTP Response
  - TCP Response
- Periodic Health Checks
  - Response Code
  - Response Time
  - Number of Retries
    - Up
    - Down



# Internal Cluster Monitoring

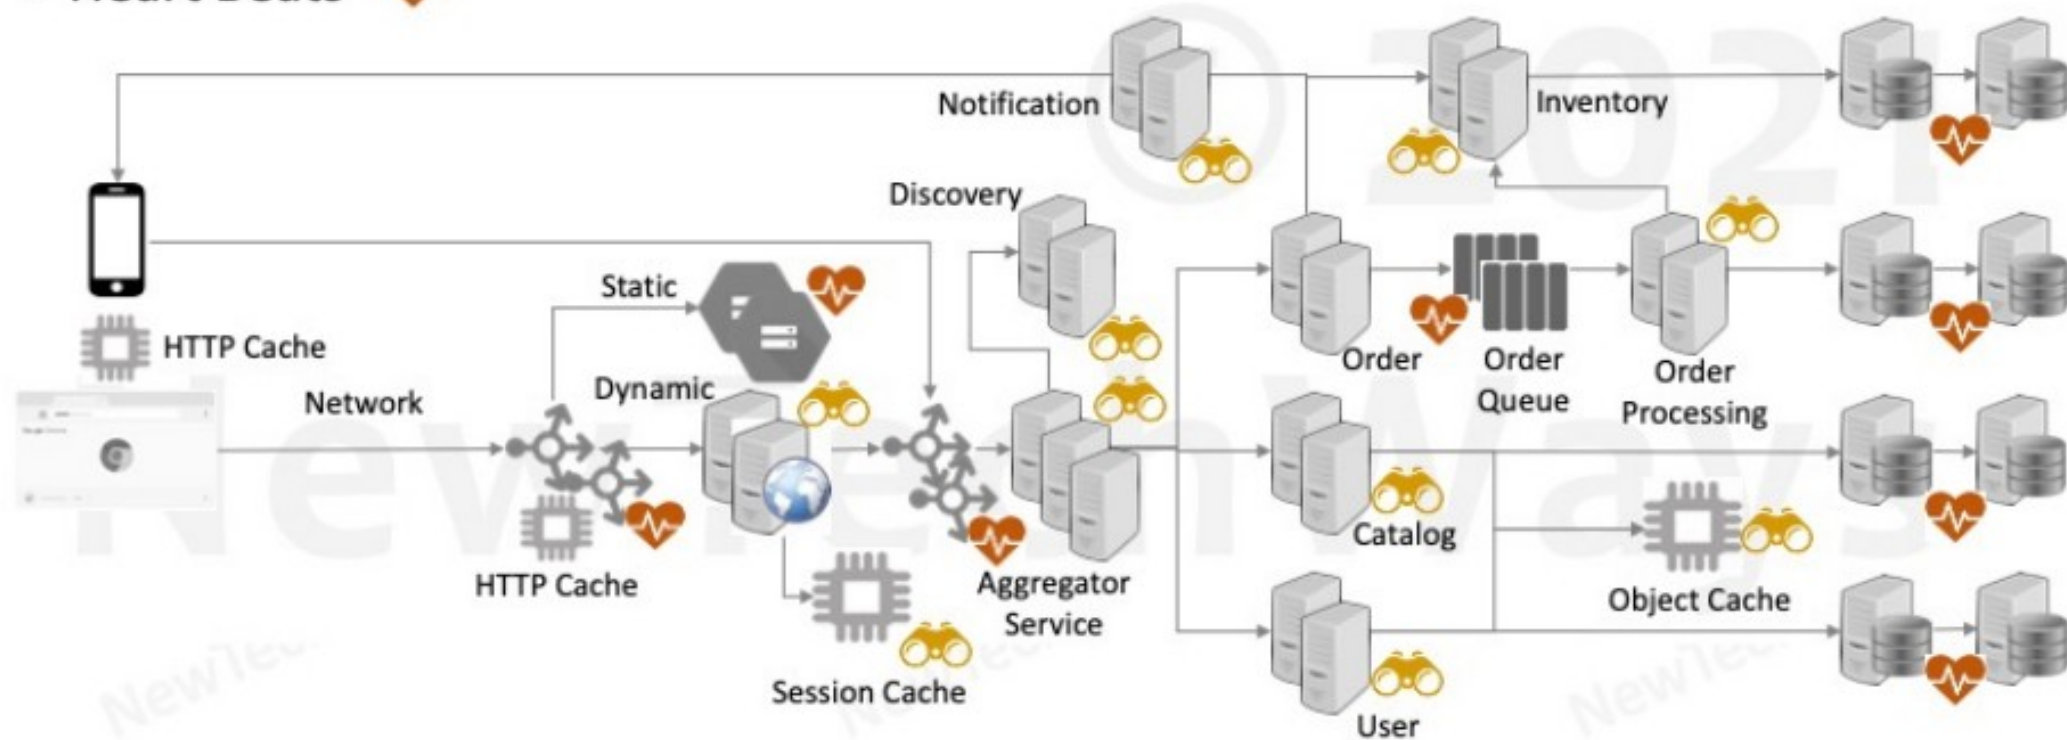
- Periodic exchange of heartbeats between redundancy cluster nodes
  - Requires protocols for communication and recovery
- Useful for stateful cluster components
  - Examples are NoSQL DB cluster and Load Balancers





# Fault Detection – Monitoring

- Health Checks 🚒
- Heart Beats ❤️

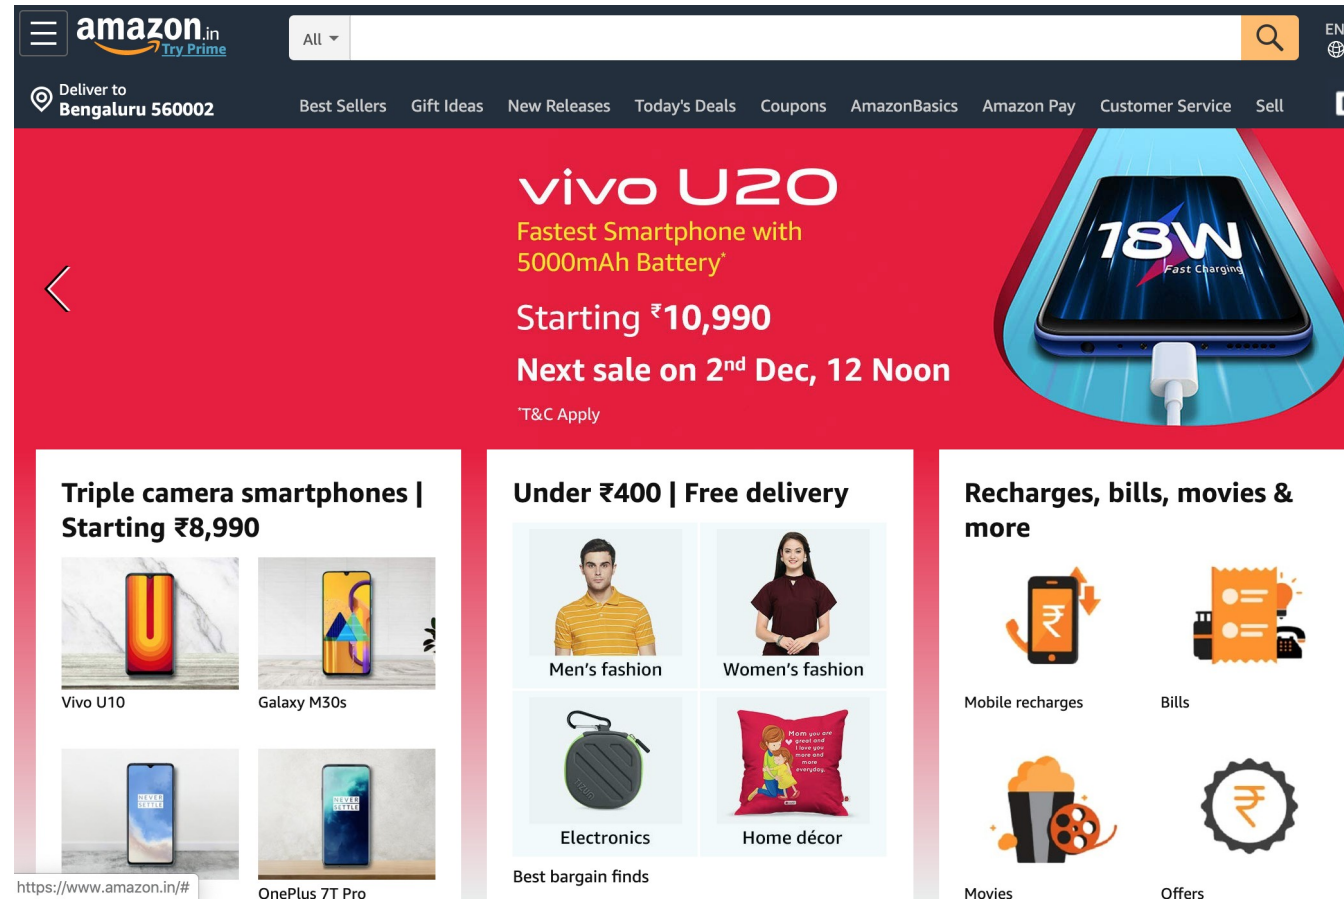




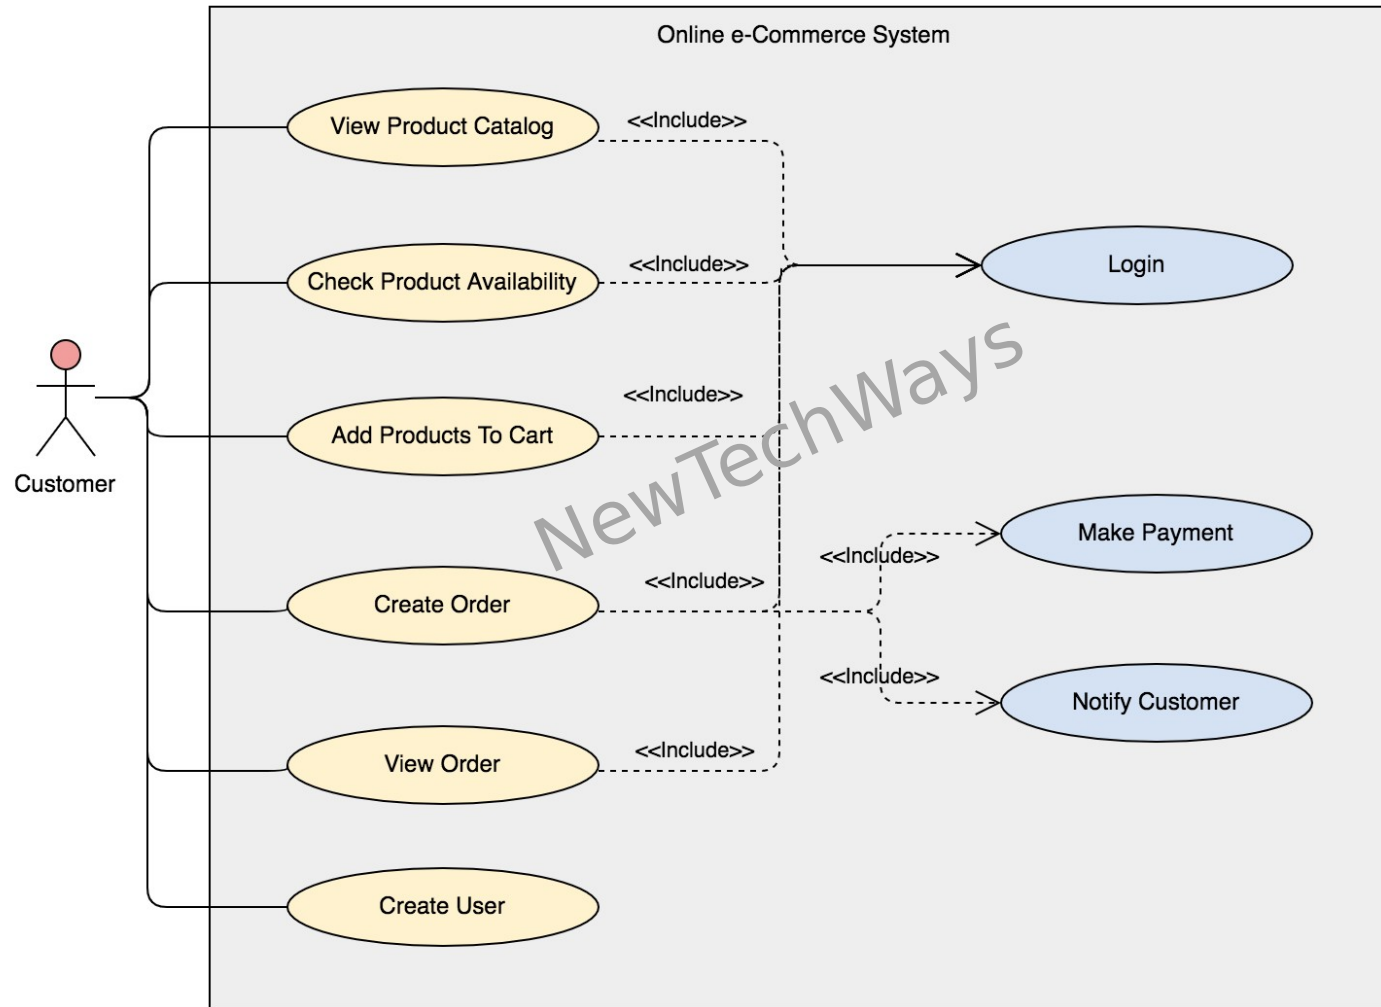
---

# Application Architecture

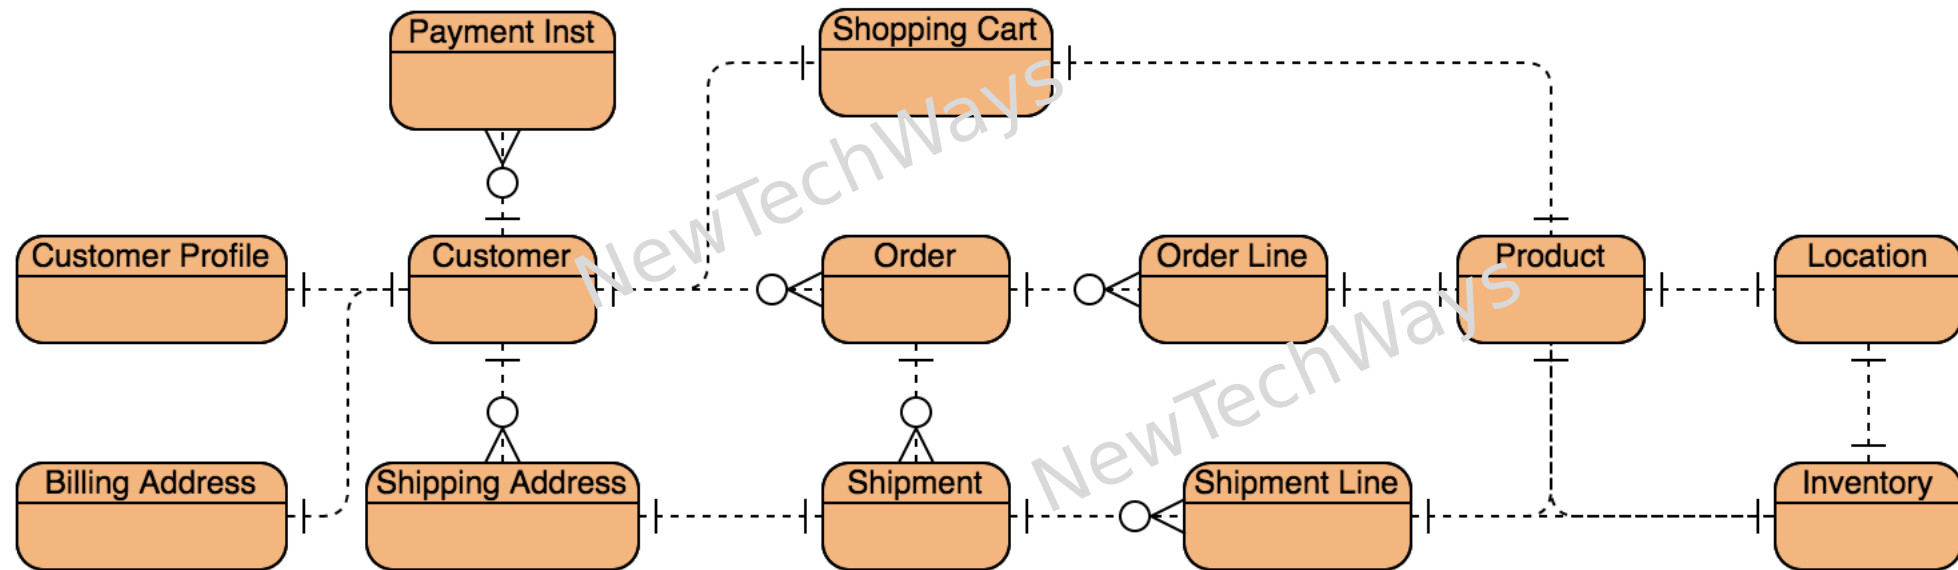
# Assume 'Architecting' a big e-commerce system



# Use Case Model




# Domain Model




# Low Level Design



GUI Prototype



Use Case Model




Sequence Diagram




Component

**Behavior**

**State**



Domain Model



Class Diagram

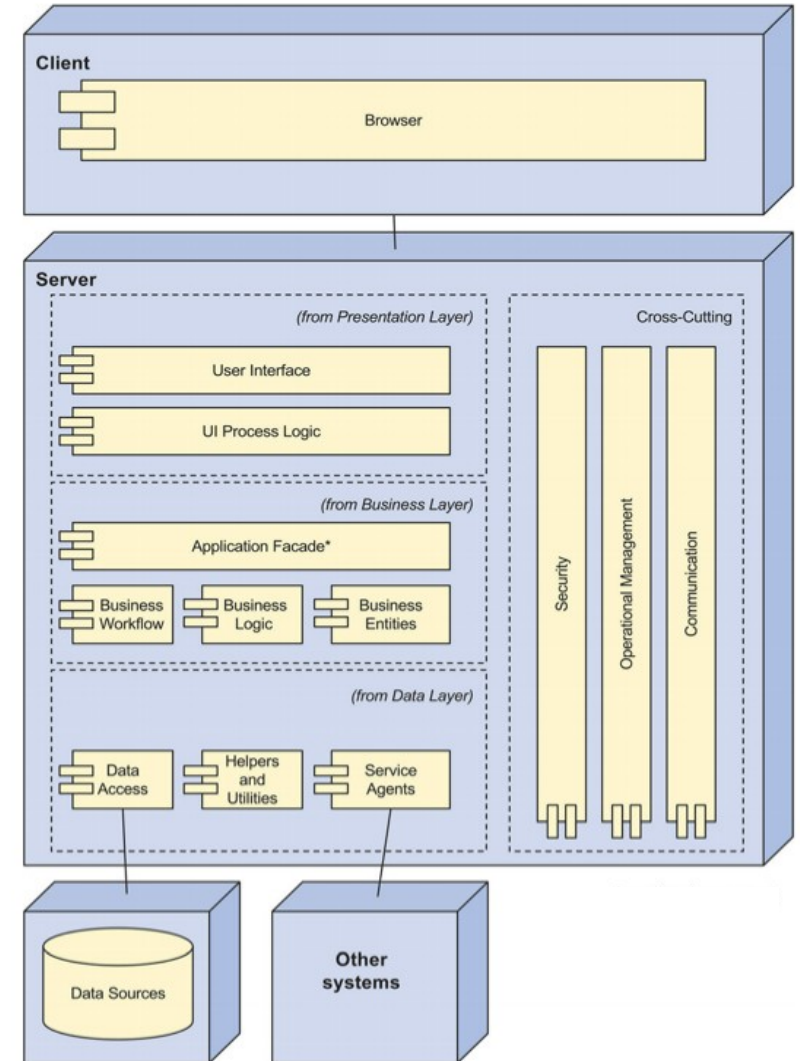



```
01100
10110
11110
```

**Code**

**Static + Transactional**

# Component Model



# What About?



High  
Availability



Response  
Latency



Global  
Customers



System &  
Data Security



File storage for  
Catalog Images



Mobile  
Support



Unstructured Data  
Storage & Analytics



Cloud  
Deployment

# Going Beyond

---



Application Architecture To System Architecture



# Application Architecture Vs System Architecture

System level challenges  
surface up in large scale systems



Scale, Reliability, Security, Deployment are  
biggest concerns for a large-scale system



# Latency Requirements

---

- 10 years ago, Amazon found that every 100ms of latency cost them 1% in sales
- Now Akamai study shows that every 100-millisecond delay in website load time can hurt sales by 6%
- Google found an extra .5 seconds in search page generation time dropped traffic by 20%

# Scalability Requirements

10 million requests/day

1K to 100K simultaneous users

100 million products

Transaction data for last 5 years

Petabytes of log data

- For 100 M products with average product description/image size of 1 MB
  - Data Storage =  $100M \times 1MB = 100 \text{ TB}$
- Logs generated and archived
  - Data Storage => in Petabytes

# Availability & Reliability Requirements

Unavailability results in

- Business loss
- Reputation loss

99.95% Availability

- Maximum cumulative disruption of 4 hours 22 minutes in a year

99.9999999999% Durability for storage systems

- Data once stored is practically never lost

Disaster Recovery

- Operations to continue even if a region goes down due to a natural calamity



# Security Requirements

---

## Infrastructure protection

- Network access
- System access
- Service access

## Data Protection

- Data sensitivity classification
- Protect data at rest
- Protect data on wire
- Data backup & replication

## Identity & Access Management

- Authentication, Authorization

---

# Designing System Architecture

# Scalability Principle

---

Monolith is an anti-pattern for Scalability

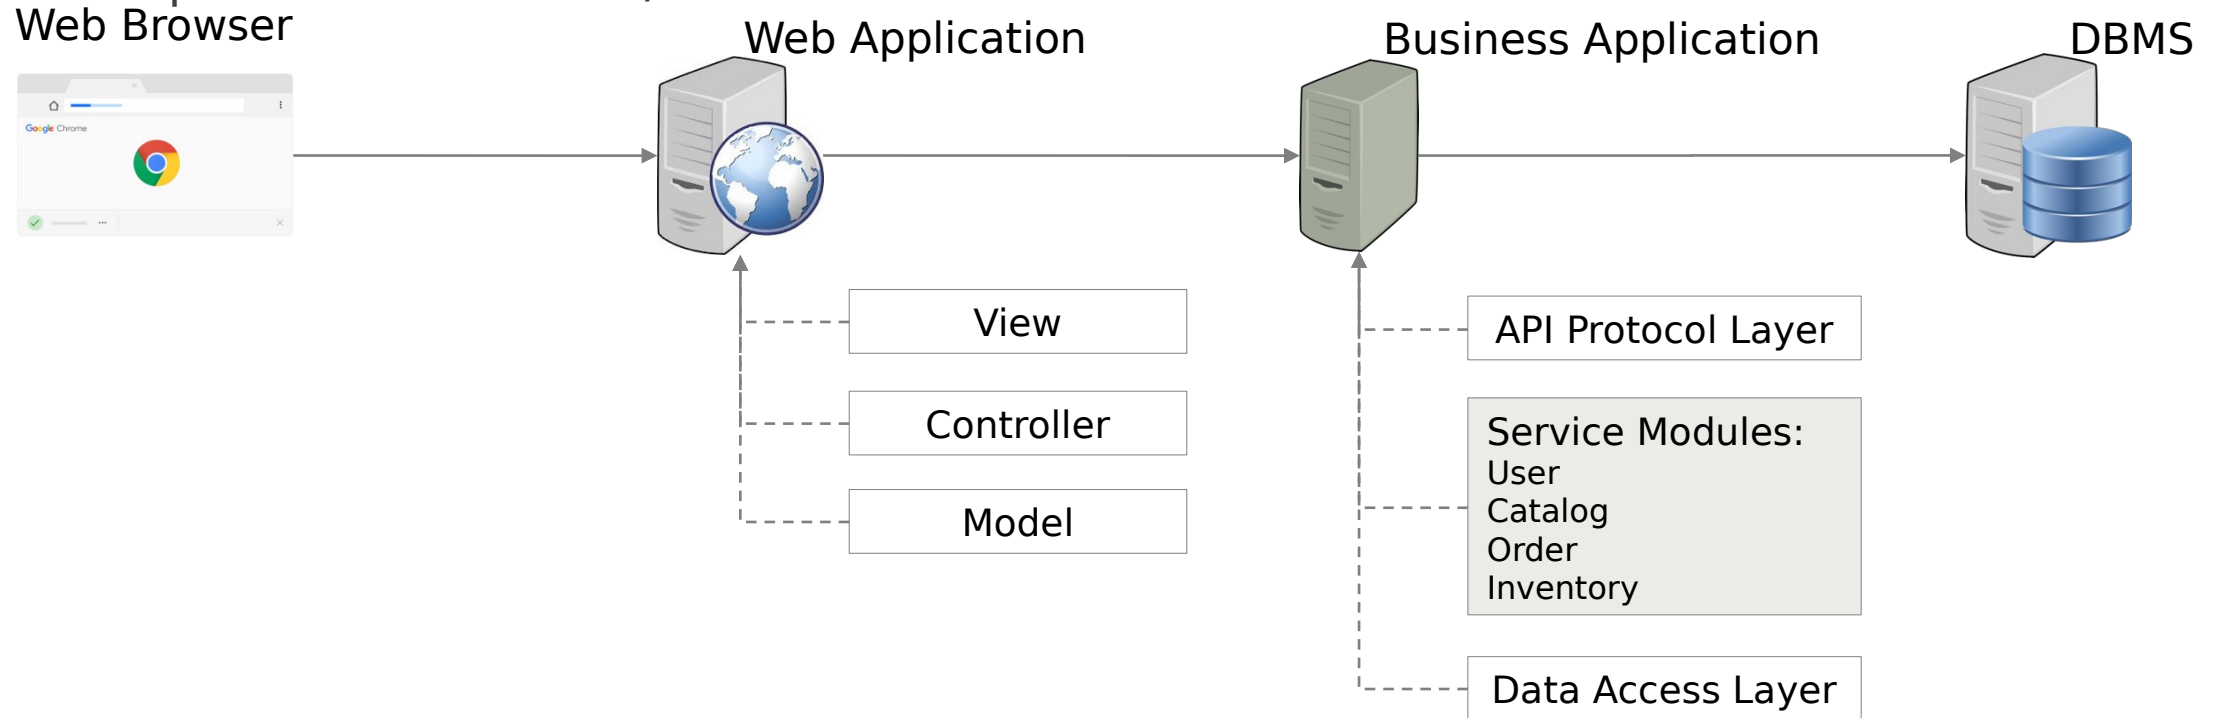
Scalability goes up with

- Decentralization
  - More specialized workers – Services
  - More workers – Instances, Processors, Threads
- Independence
  - Multiple workers are as good as a single worker if they can't work independently
  - They must work concurrently to maximum extent

# Modularity

Scalable architecture starts with modularity


- Provides the foundation for breaking a system function/service into more specialized functions/services





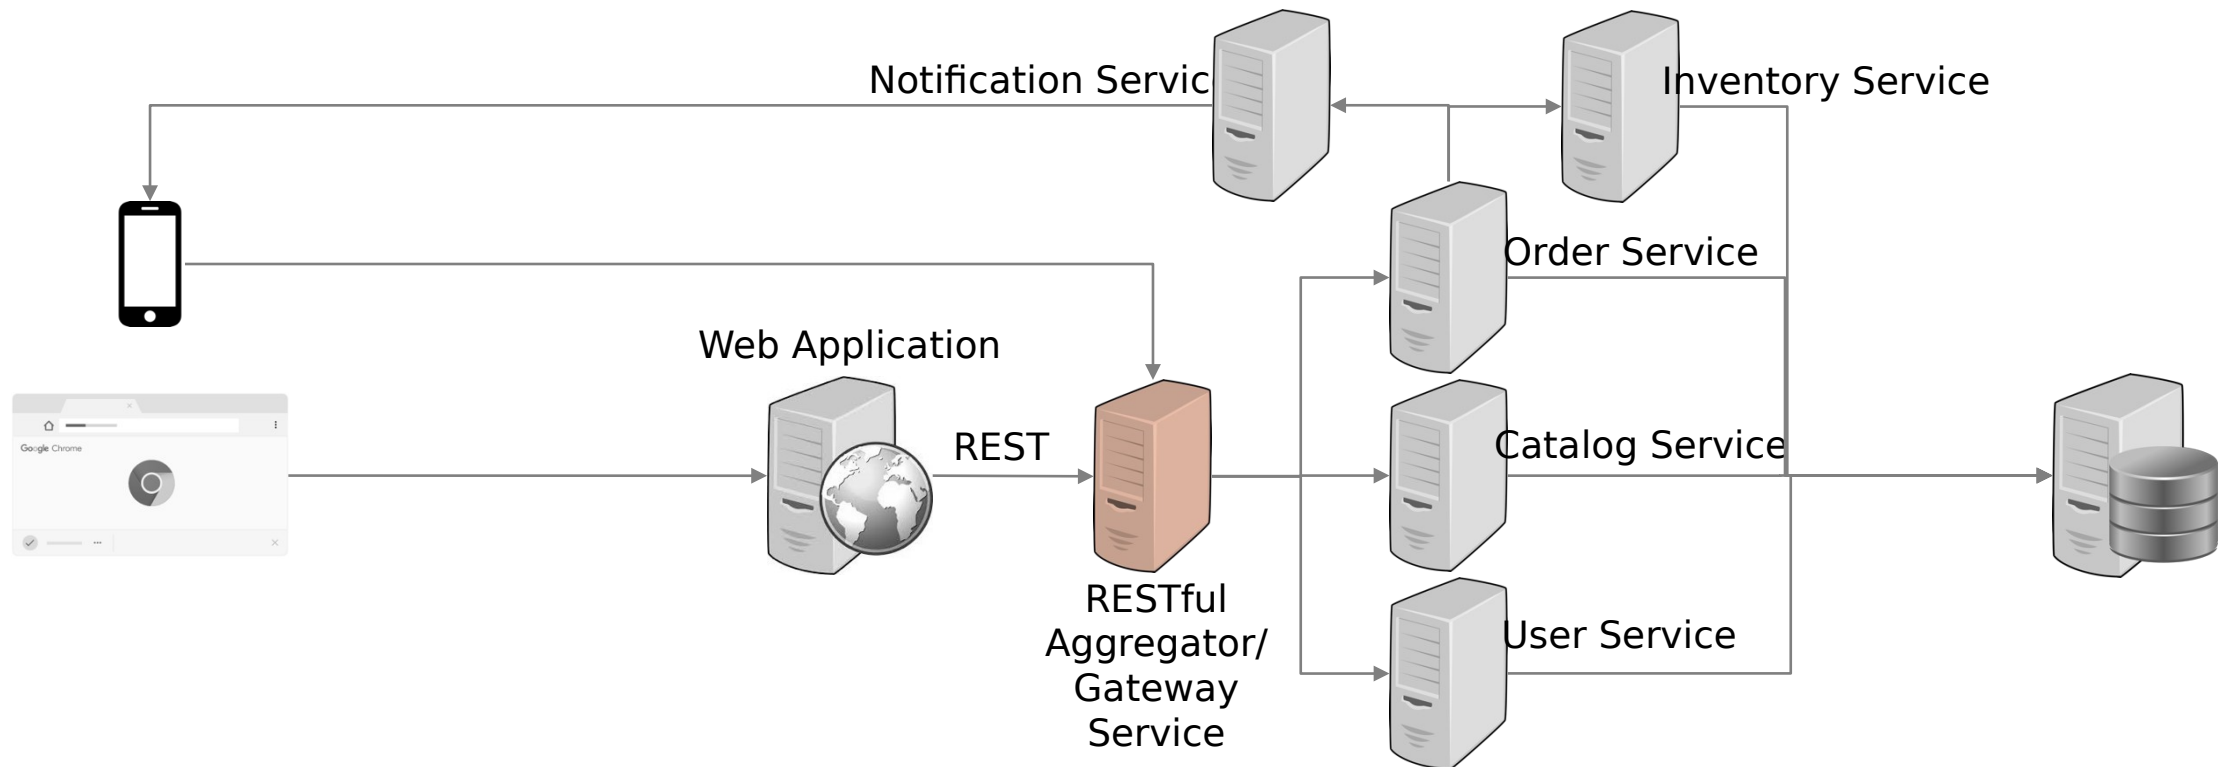
# Specialized Services – WebServices

Services can be scaled differently e.g., Number of instances



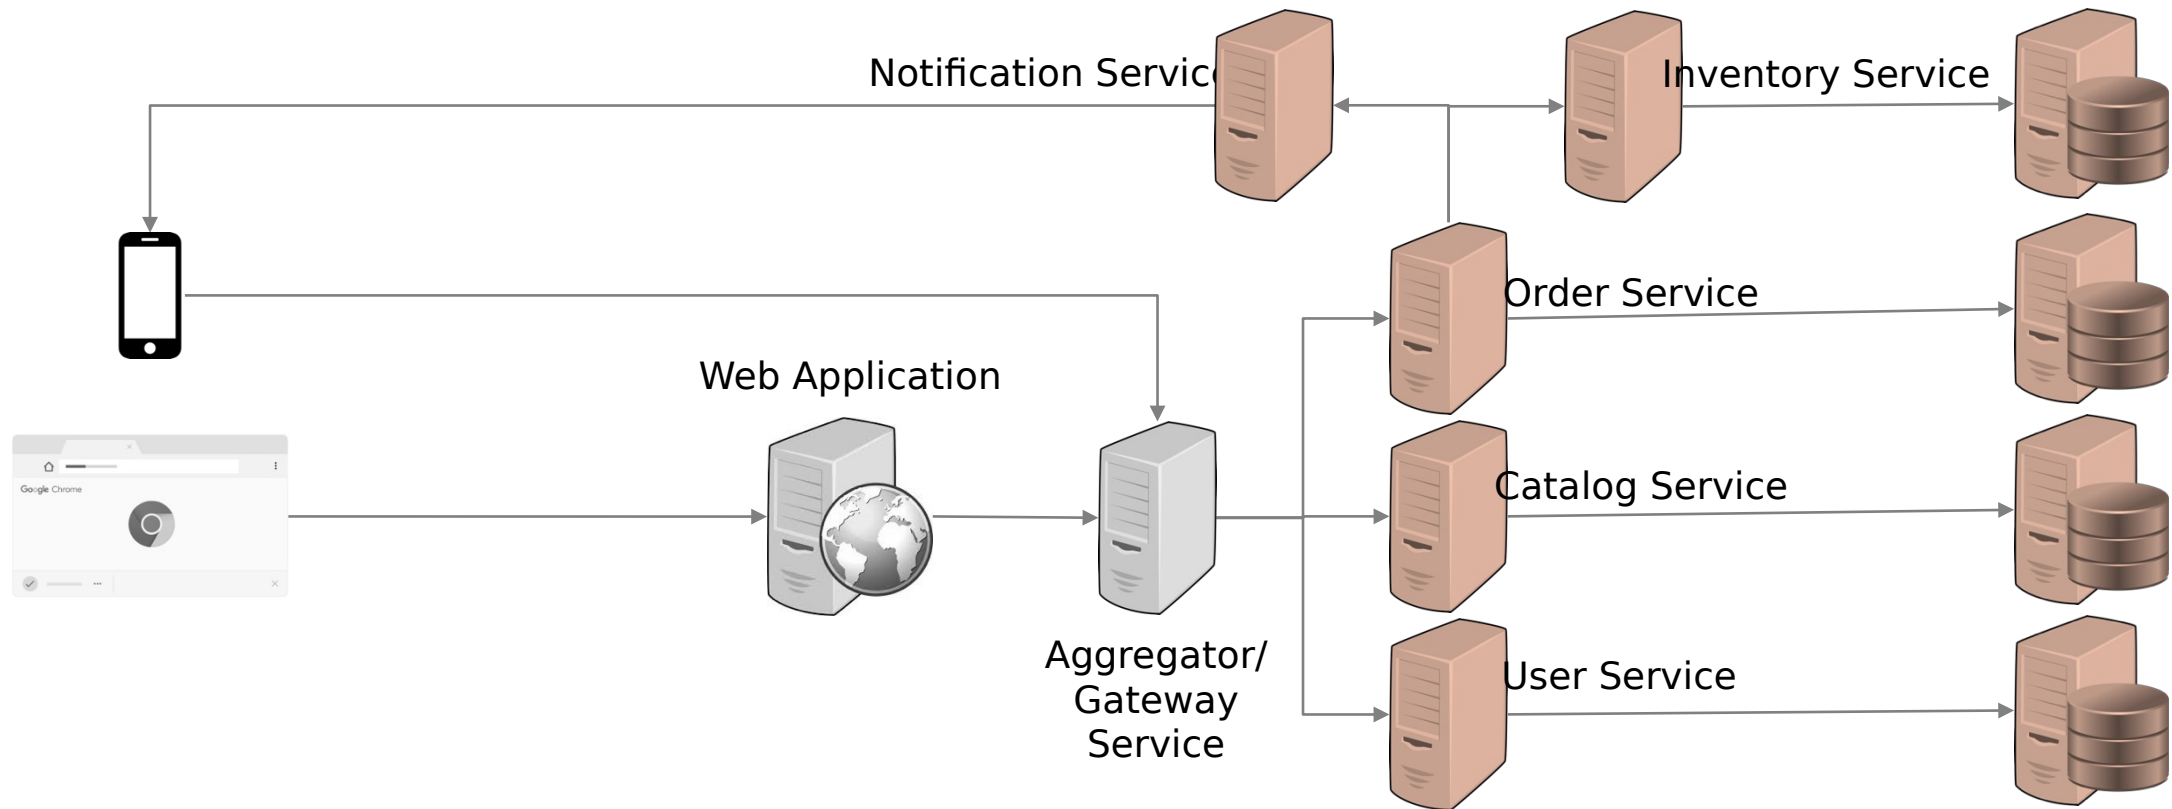
# Aggregator Service & RESTful API – Mobile Support

REST for external interface interoperability & Mobile Support



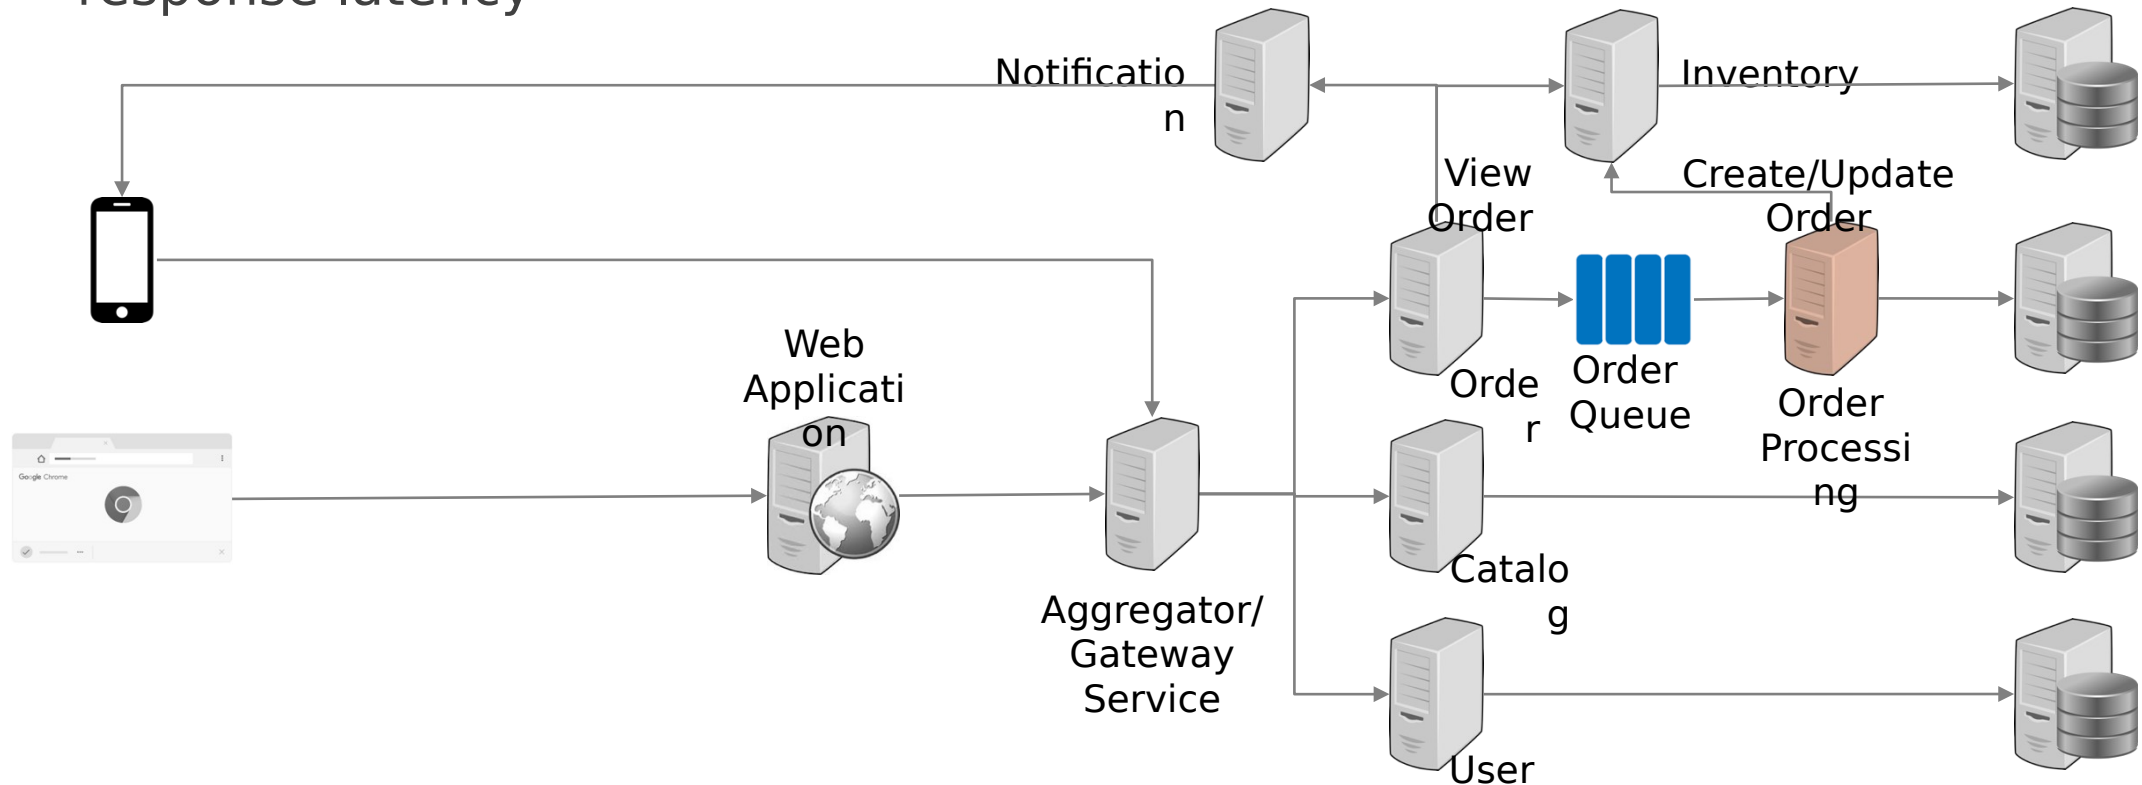
# Independent Services – MicroServices

Micro-Services can be scaled differently and deployed independently



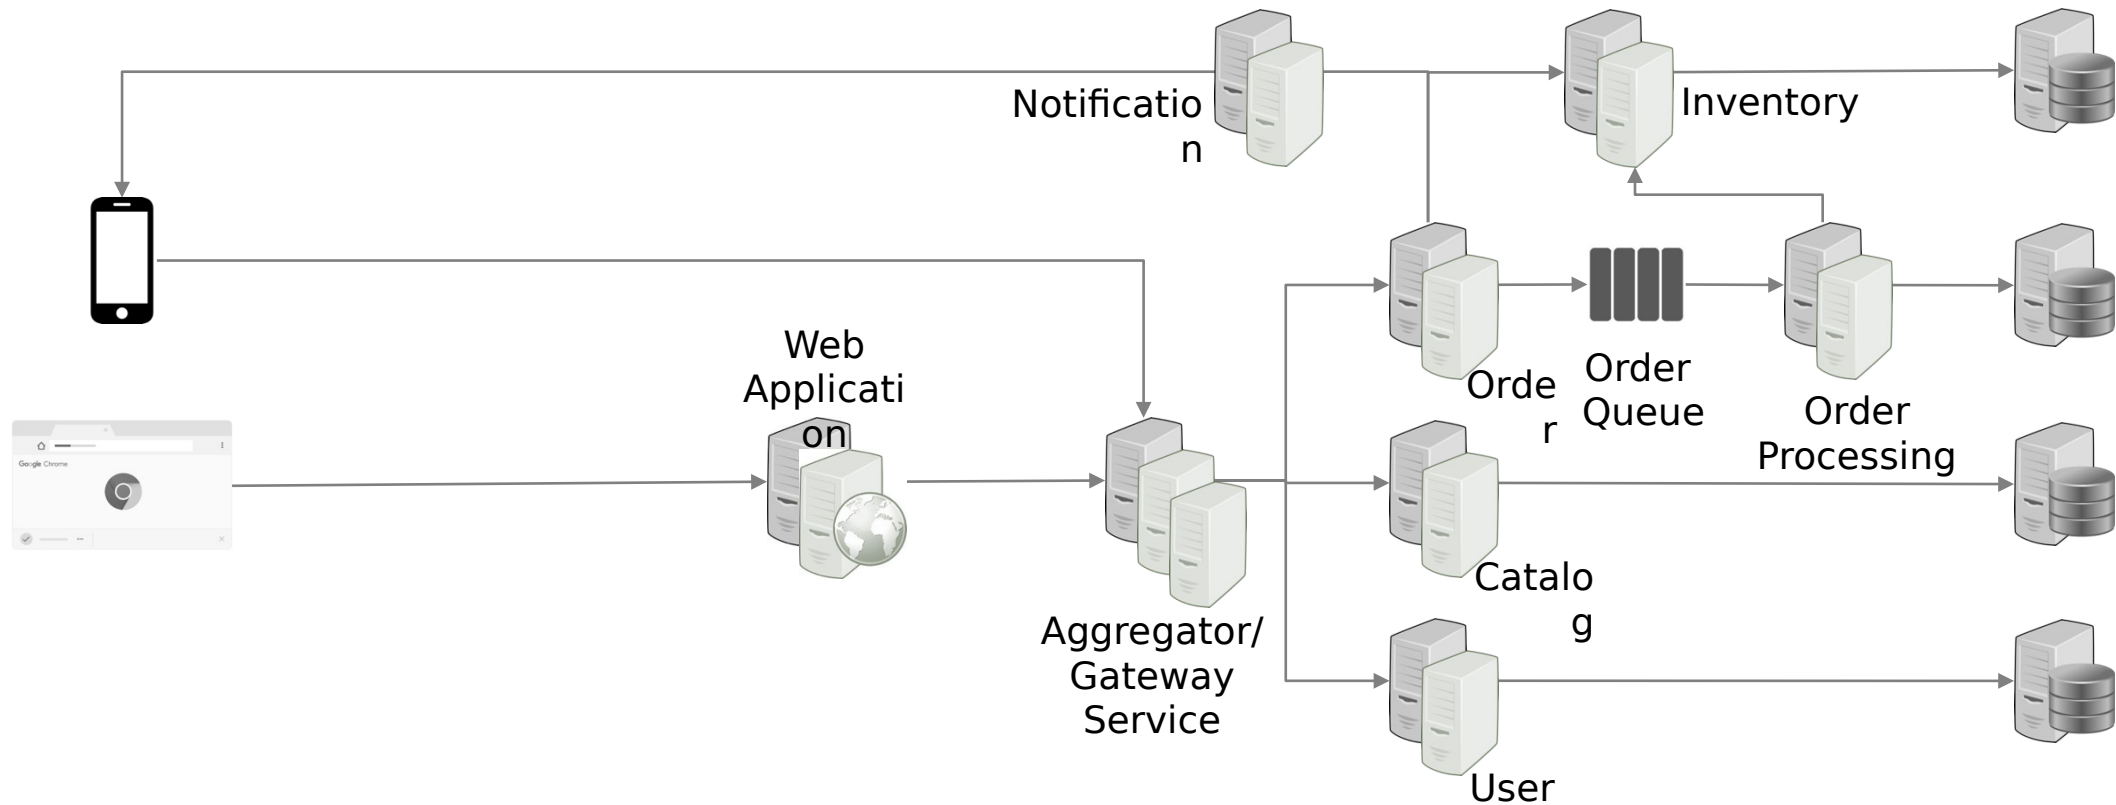
# Asynchronous Services

Updates can be done asynchronously to buffer peak loads and to reduce response latency



# Stateless Replication

Replication provides more computation power



# Reliability Principle

---

## Reliability

- Normal functioning even in the presence of faults

## Availability

- Always available even in the presence of faults

Reliability and Availability are achieved mainly through Fault Tolerance

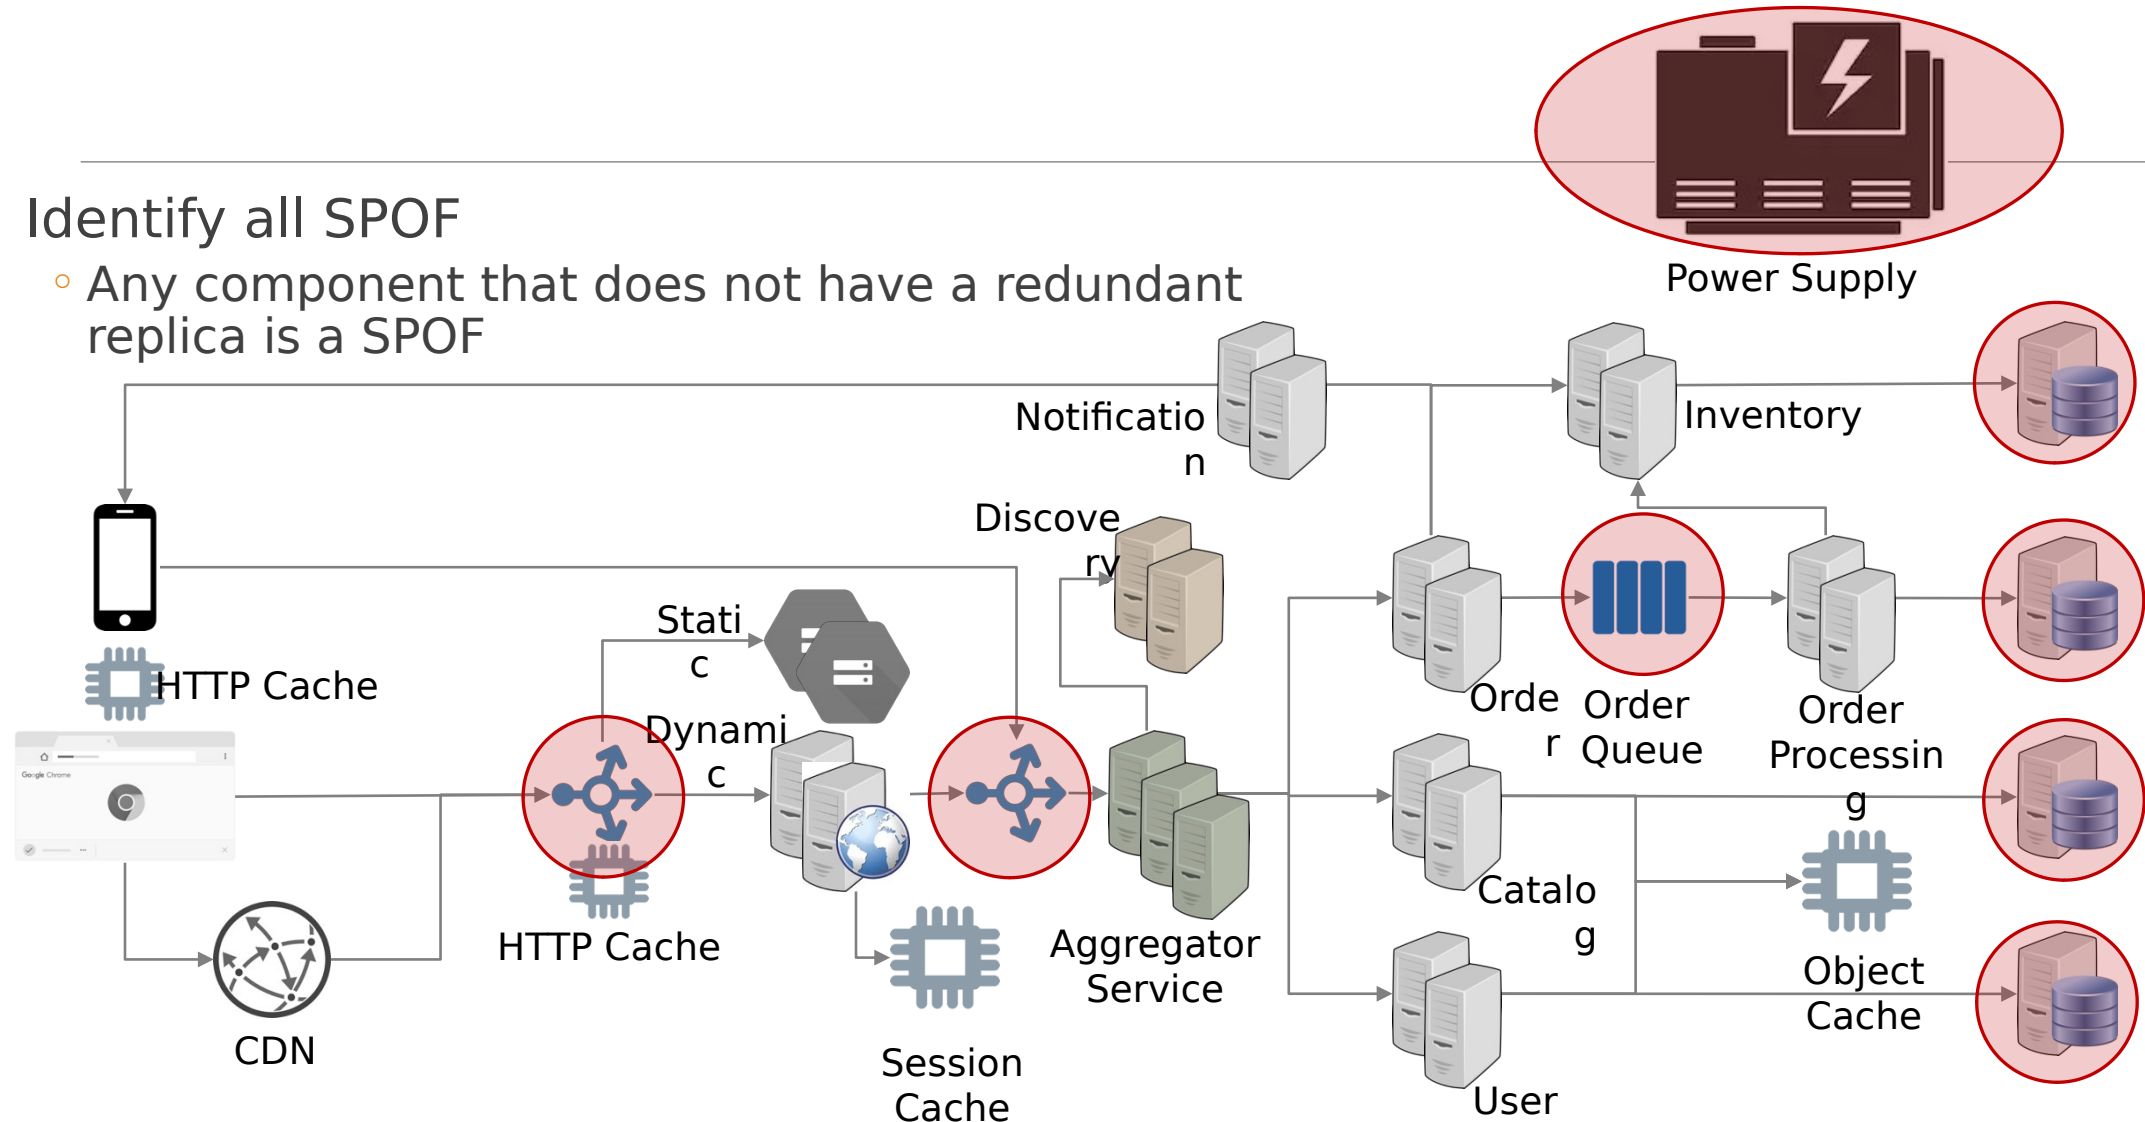
## Fault Tolerance requires

- Provisioning Redundancy
- Fault Detection Mechanisms
  - Health-checks, heart-beats

# Single Point Of Failures

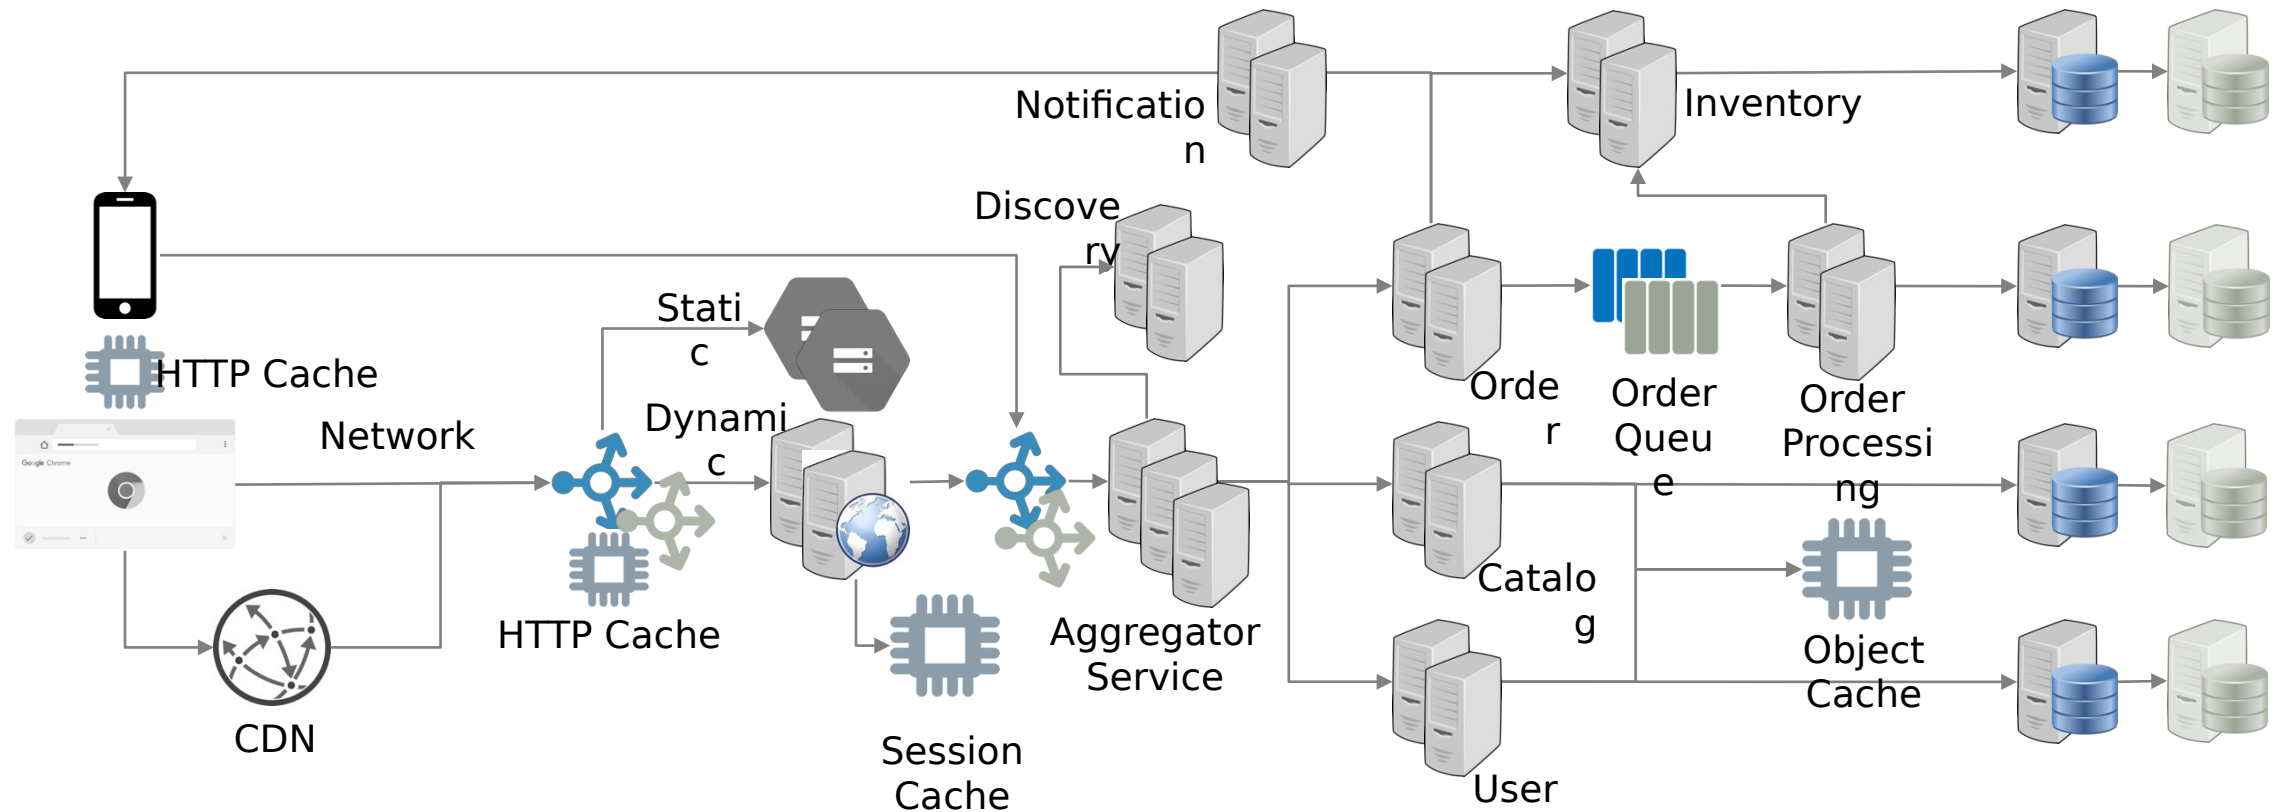

## Identify all SPOF

- Any component that does not have a redundant replica is a SPOF



# Redundancy

Components that are SPOF requires redundancy

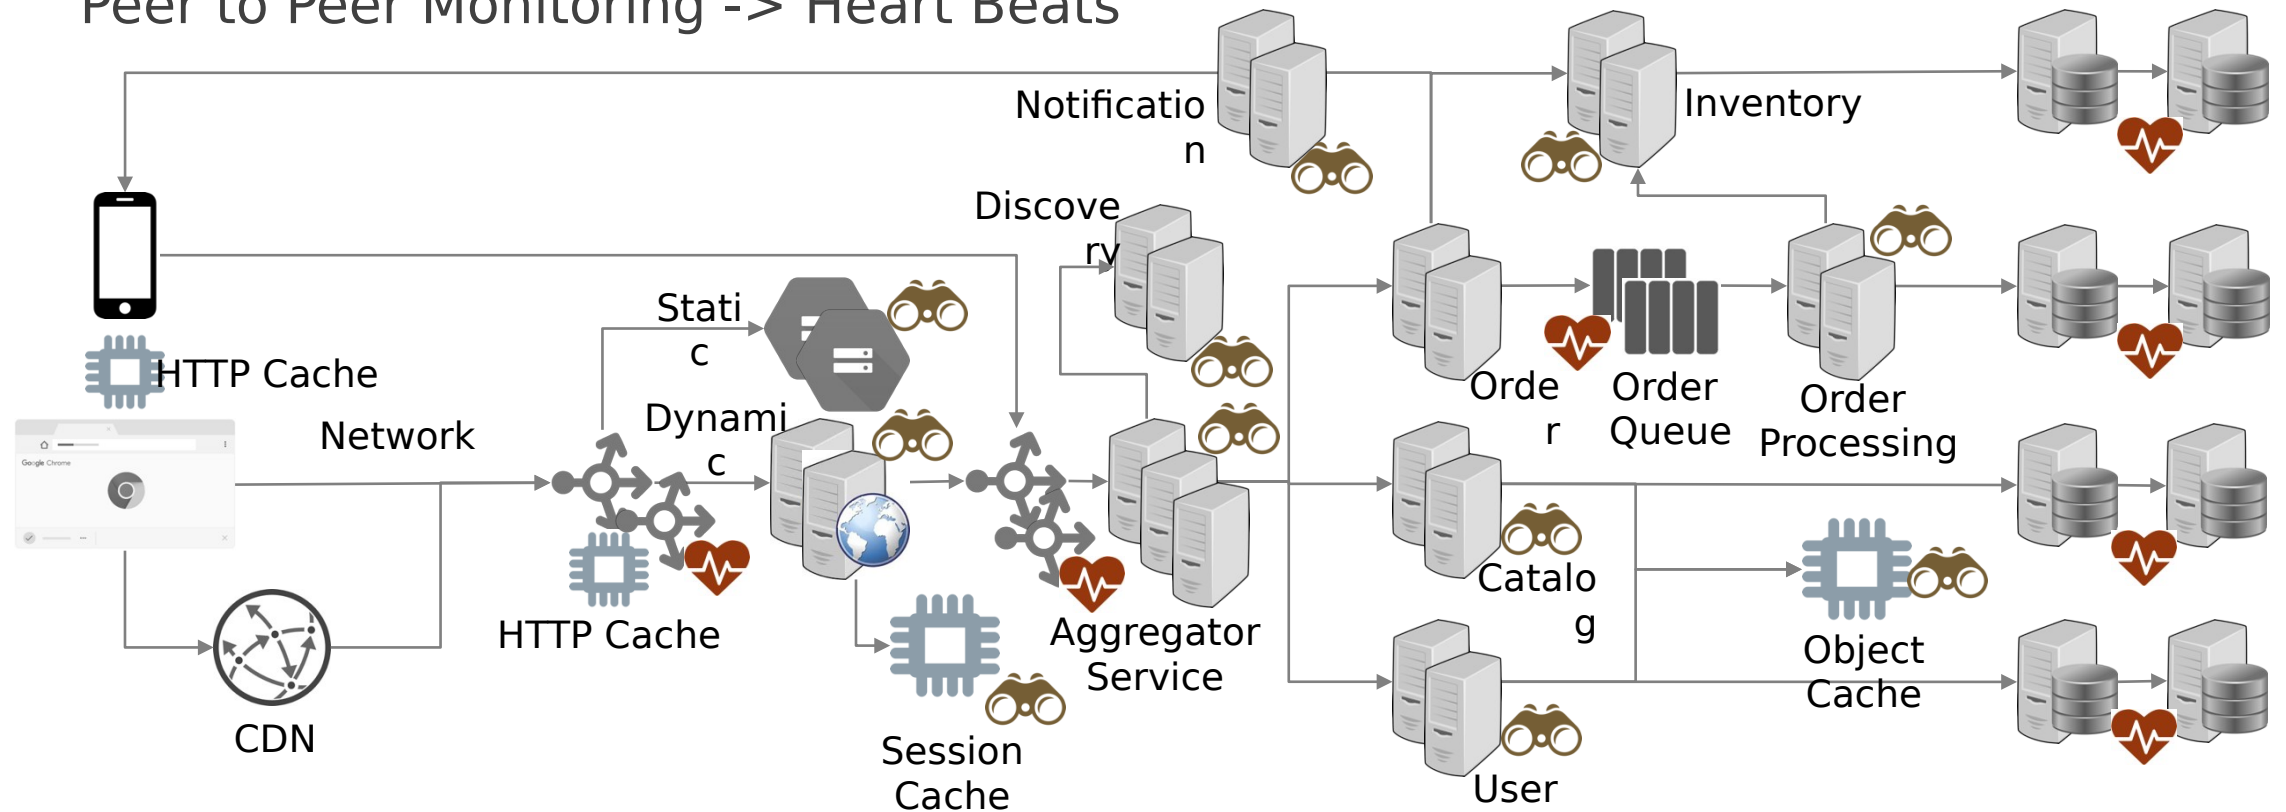





# Monitoring For Fault Detection


Hierarchical Monitoring -> Health Checks 

Peer to Peer Monitoring -> Heart Beats



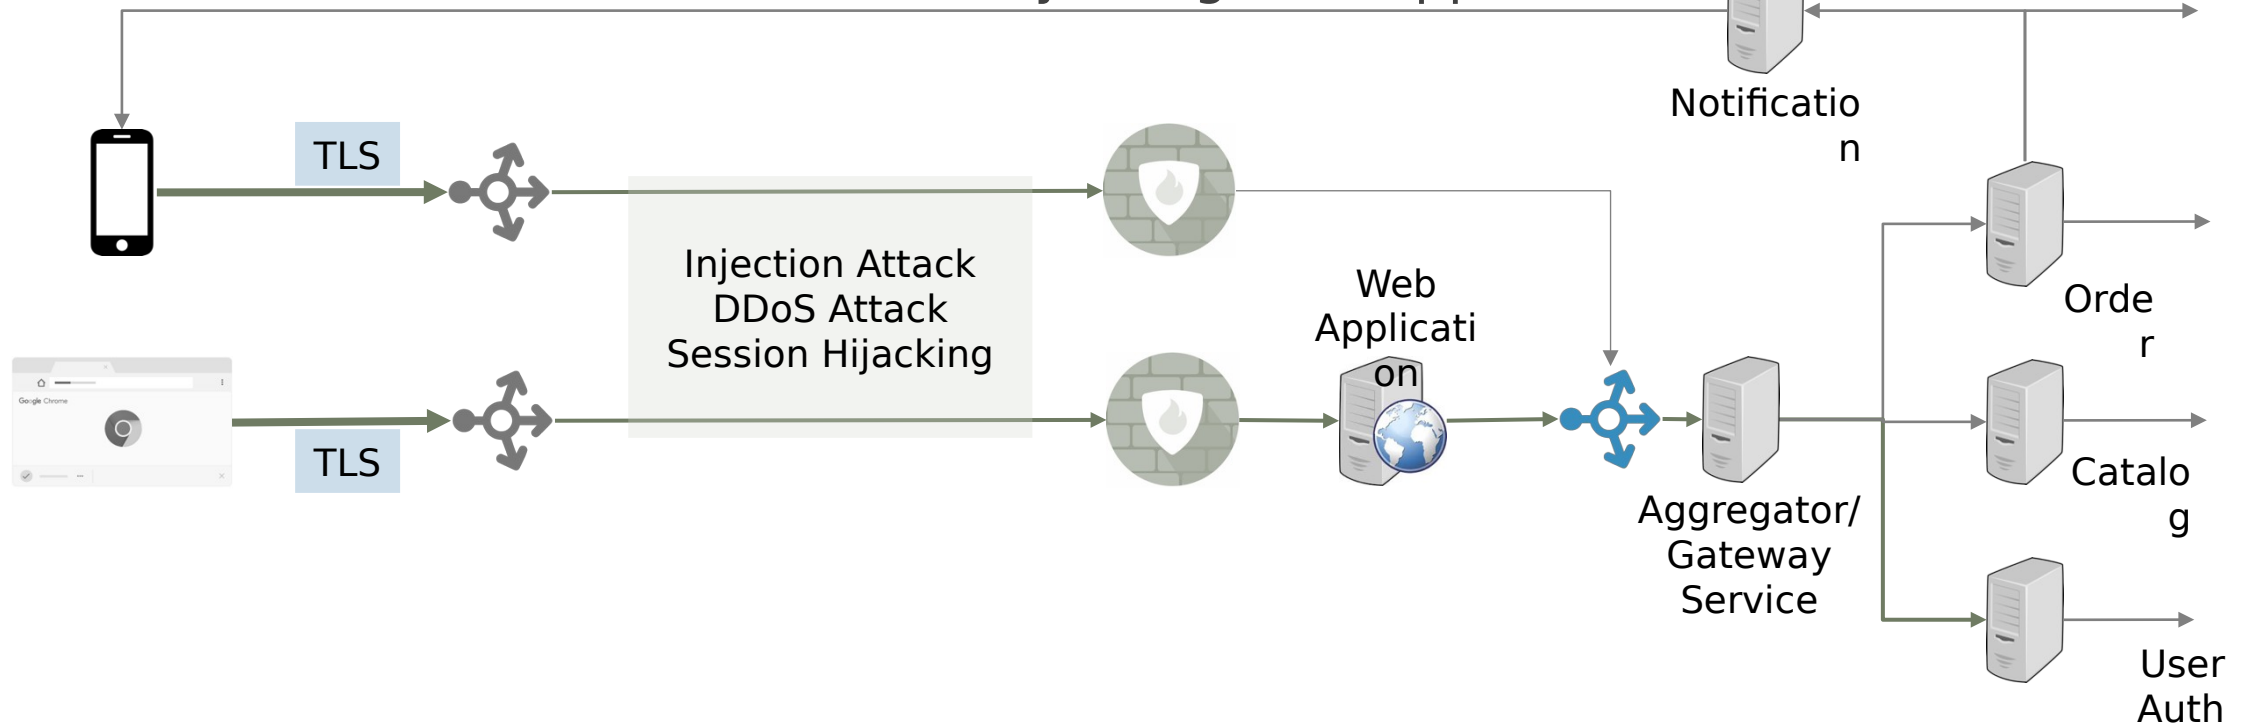
# Zonal Redundancy

Zonal redundancy for fault isolation



# Security-Web Application Firewalls

Multiple security related attacks can be prevented by inspecting requests for common attacks and vulnerabilities by using Web Application Firewalls





**HAVE A GOOD DAY!**